

Using KVM as a Transparent Hardware Abstraction Layer

David Ahern



Agenda

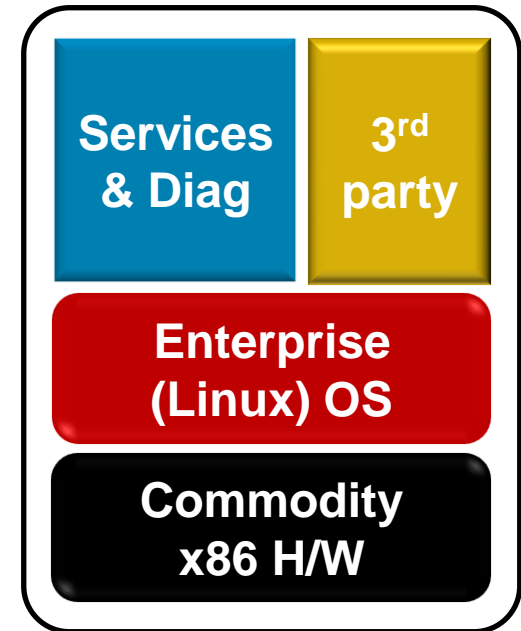
- Introduction
 - Appliance-like, enterprise products
 - Component lifecycles – and impacts of their misalignments
- Design Overview
 - A high level look at how KVM is used as a solution for handling the impacts of the short lifecycle commodity H/W
- Design Details
 - A closer look at details such as installation, networking, monitoring, security and performance
- Final Thoughts

Problem Introduction



Appliance-Like, Enterprise Products

- “Bundled” data center solution
 - Product tied to H/W and includes OS
- Closed-box design
 - Product handles firmware & BIOS updates, configures H/W (RAID & BIOS), and includes H/W in monitoring and serviceability
- Commodity x86 Hardware
 - Typical data center servers using Xeon-class processors: e.g., IBM x3250 & x3650, HP DL320 & DL380
- Enterprise Linux OS



Product Deployment

- Customer purchases specific version of product
 - Testing and evaluations (e.g., features & security) in lab environment before going into production environment
 - Time investment with a product version
- Preference to stay on version for several years
 - Significant impact to customer to change versions
 - Minor dot release updates ***might*** be acceptable
- Need for new servers over lifetime of deployment
 - Expansion - add more nodes to expand capacity or increase fault tolerance
 - Server replacement in the case of hardware failures
- Original H/W model deployed may be past end-of-sale
 - Need support for newer H/W generations in “legacy” releases

Independent Component Lifecycles

- Product with desired support for 5 years



- Enterprise OS with support lifecycle of 7 years



- New hardware enablement limited to subset of lifetime
- Commodity x86 hardware revisions every 2-3 years



- Support window through OS lifetime, with N, N-1 preference (e.g., support for RHEL4 and RHEL5)

Lifecycle Misalignments

- Lifecycles of the H/W, OS and product will never align
 - Independent vendors, setting timelines based on their goals
- H/W changes cause the most pain
 - New H/W has to be added to product release which means re-spin, maintenance or dot release
 - New H/W typically requires an OS update
- OS updates introduce churn to the product
 - even “minor” ones inevitably have negative impacts
- Hardware and OS are “infrastructure”
 - Changes are essentially cost overhead for product, i.e., profit impacting

Handling H/W Turnover Within Release

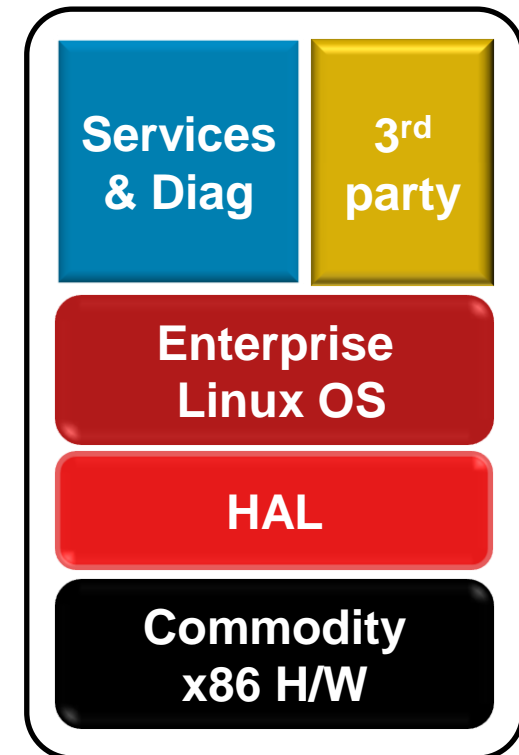
- Lifetime buys of servers
 - Requires guess-timate on sales which will never be right
- Backport OS required for H/W version
 - Non-trivial amount of grunt work to change OS versions, especially if a major version change is required
- Reduce Product lifecycle
 - Shorten sales/support periods
 - Creates unhappy customers
- KVM as a transparent HAL
 - Separates H/W from revenue generating services
 - Use technology to address problem

Design Overview



Separating Product and H/W


- Leverage virtualization technology to provide an abstraction layer
 - HAL handles H/W
 - OS for revenue generating services sees compatible virtual H/W
- Isolates primary services from hardware
 - HAL is a translator between product OS and H/W
 - H/W and the HAL can change with less churn to the overall product



Hardware Abstraction Layer

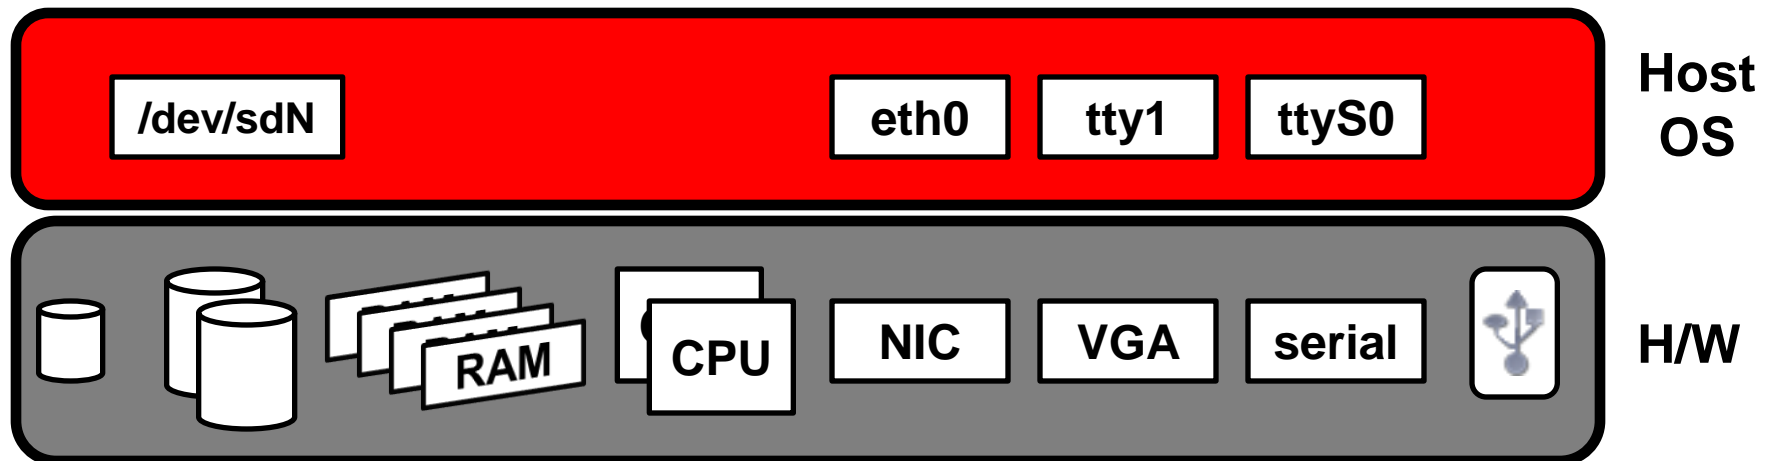
- Just another layer below product OS
 - Like a BIOS or firmware
- HAL is a newer version of a Linux OS
 - Compatible with new H/W generation
 - e.g., RHEL5 for HAL, RHEL 3 or 4 for product OS
- HAL == Host OS
 - Use virtualization to run product
 - Focused 1:1 virtualization deployment

Start with Commodity Server



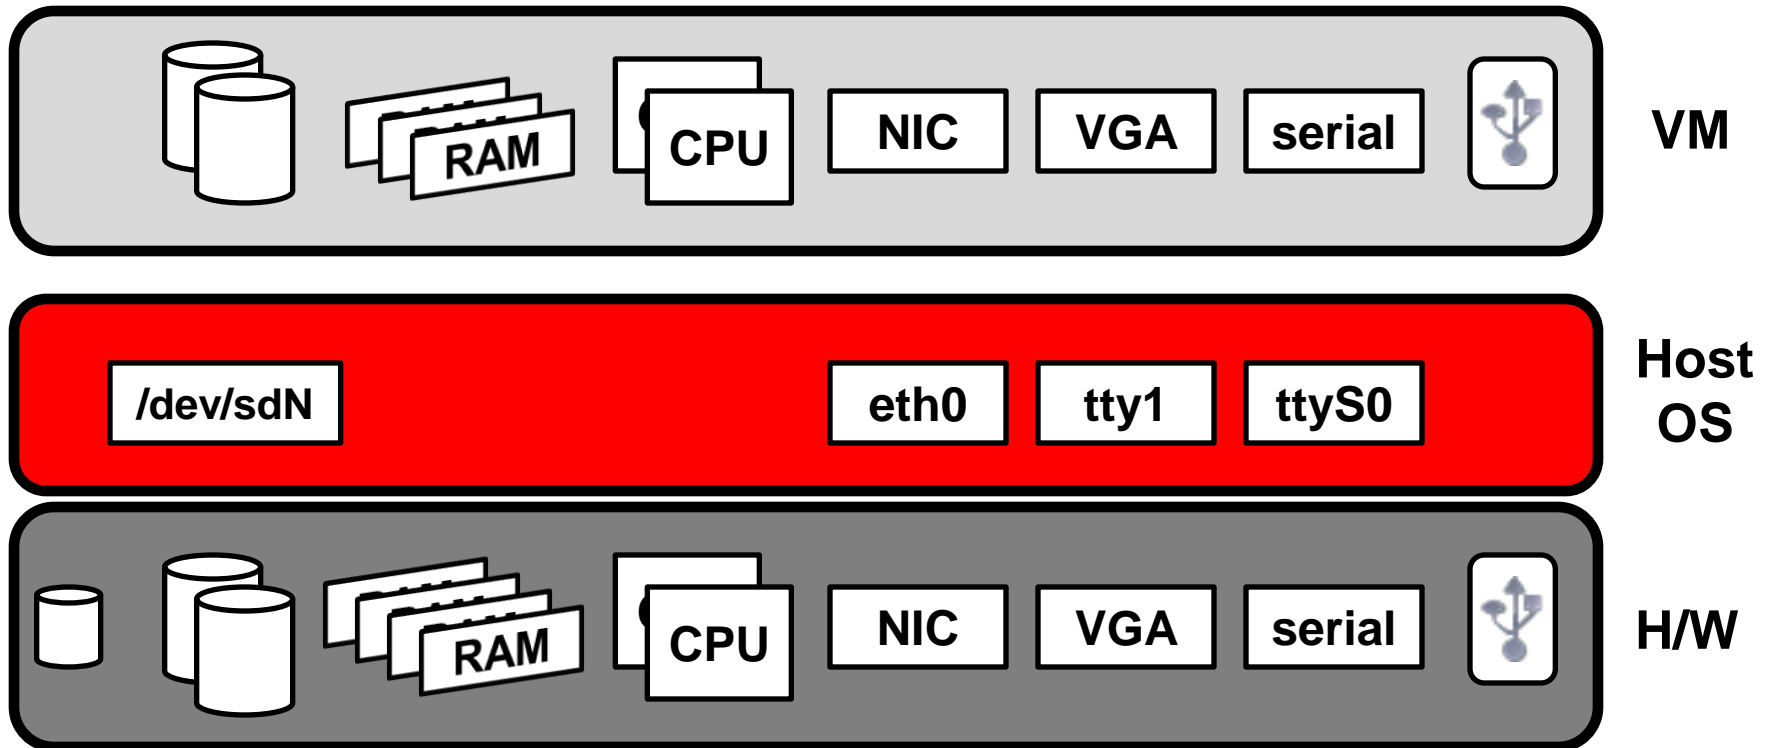
Install Host OS

- Host OS provides drivers to control real H/W



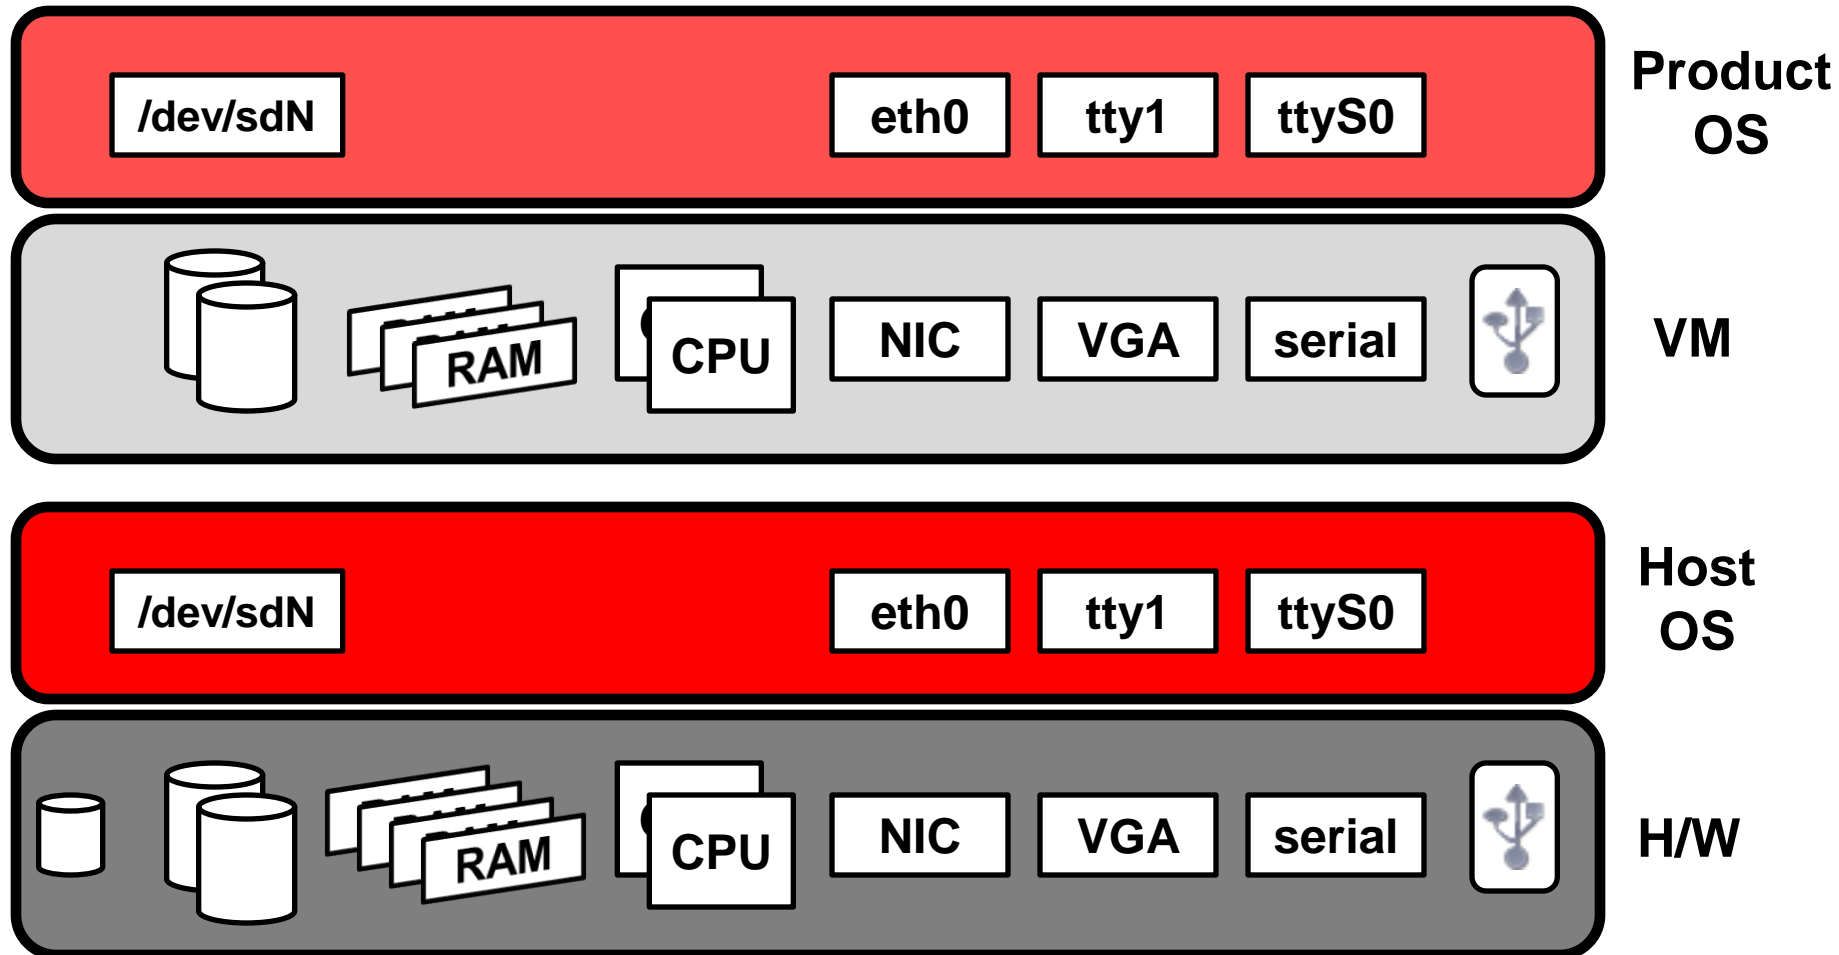
Start Virtual Machine

- Virtual Machine created with H/W “allocation”



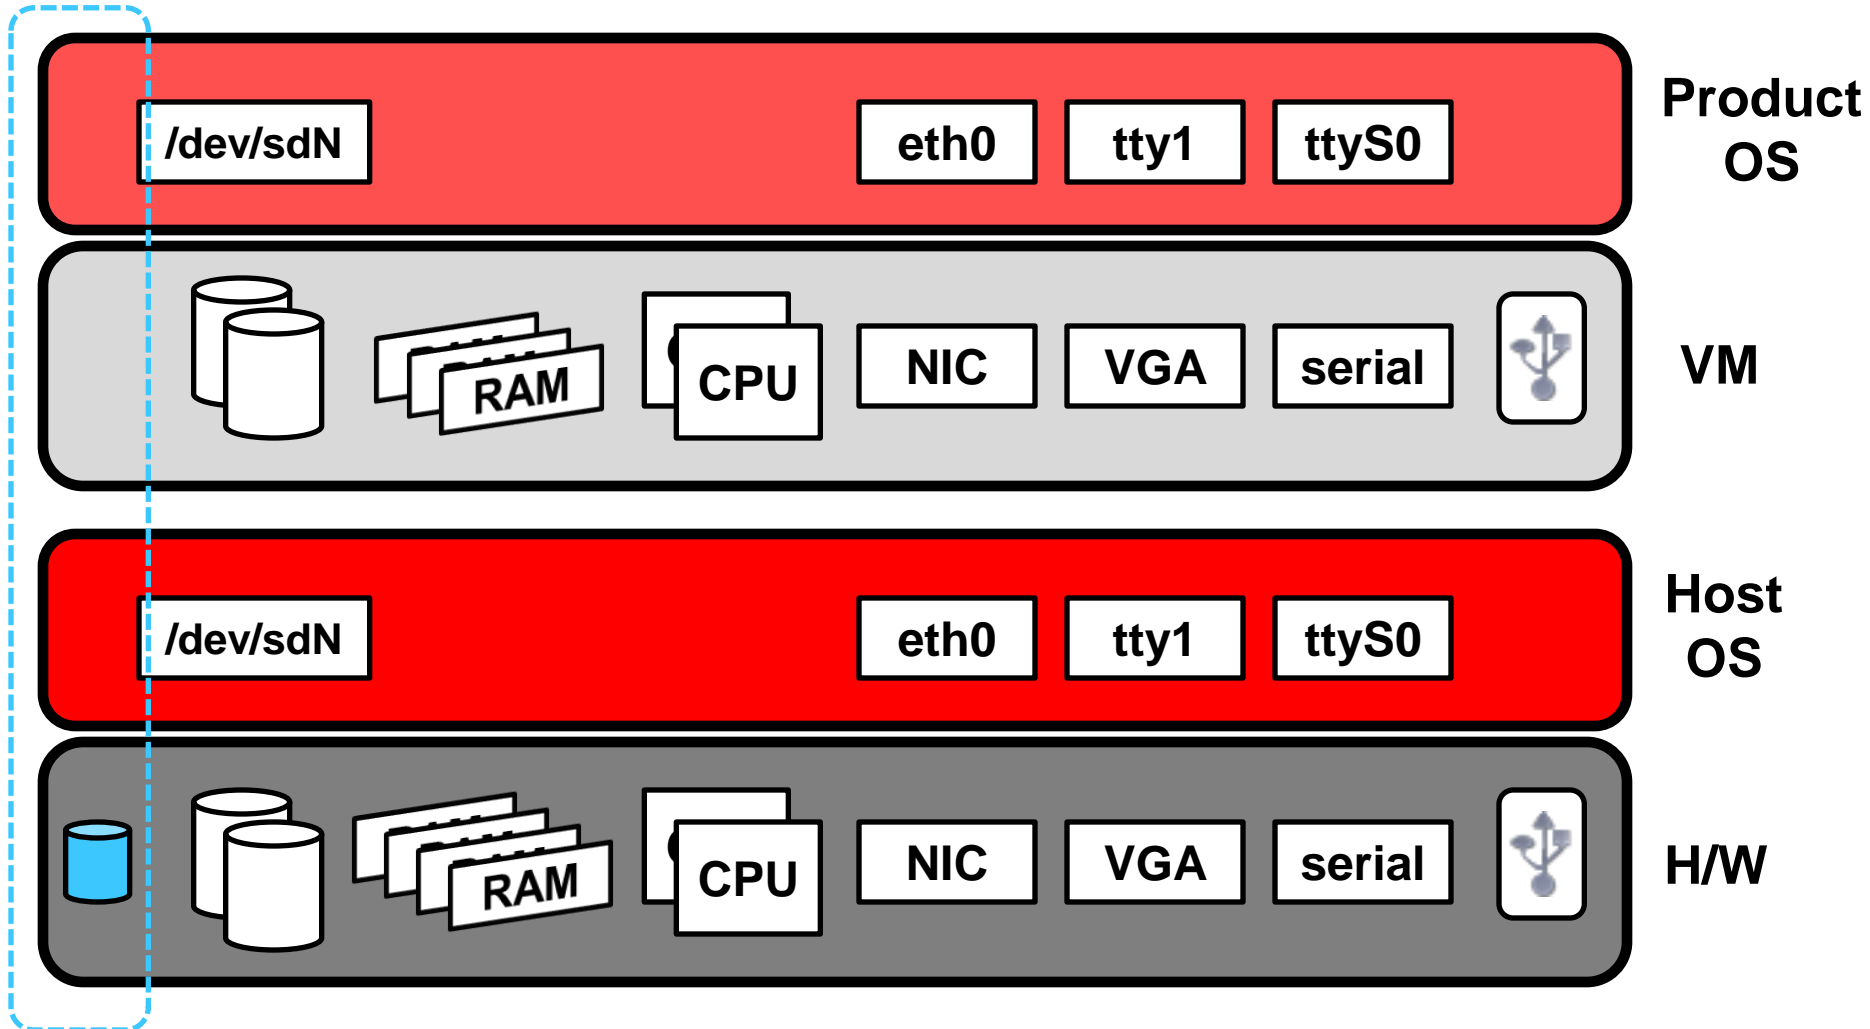
Install Product OS and Services

- Product OS is compatible with VM H/W



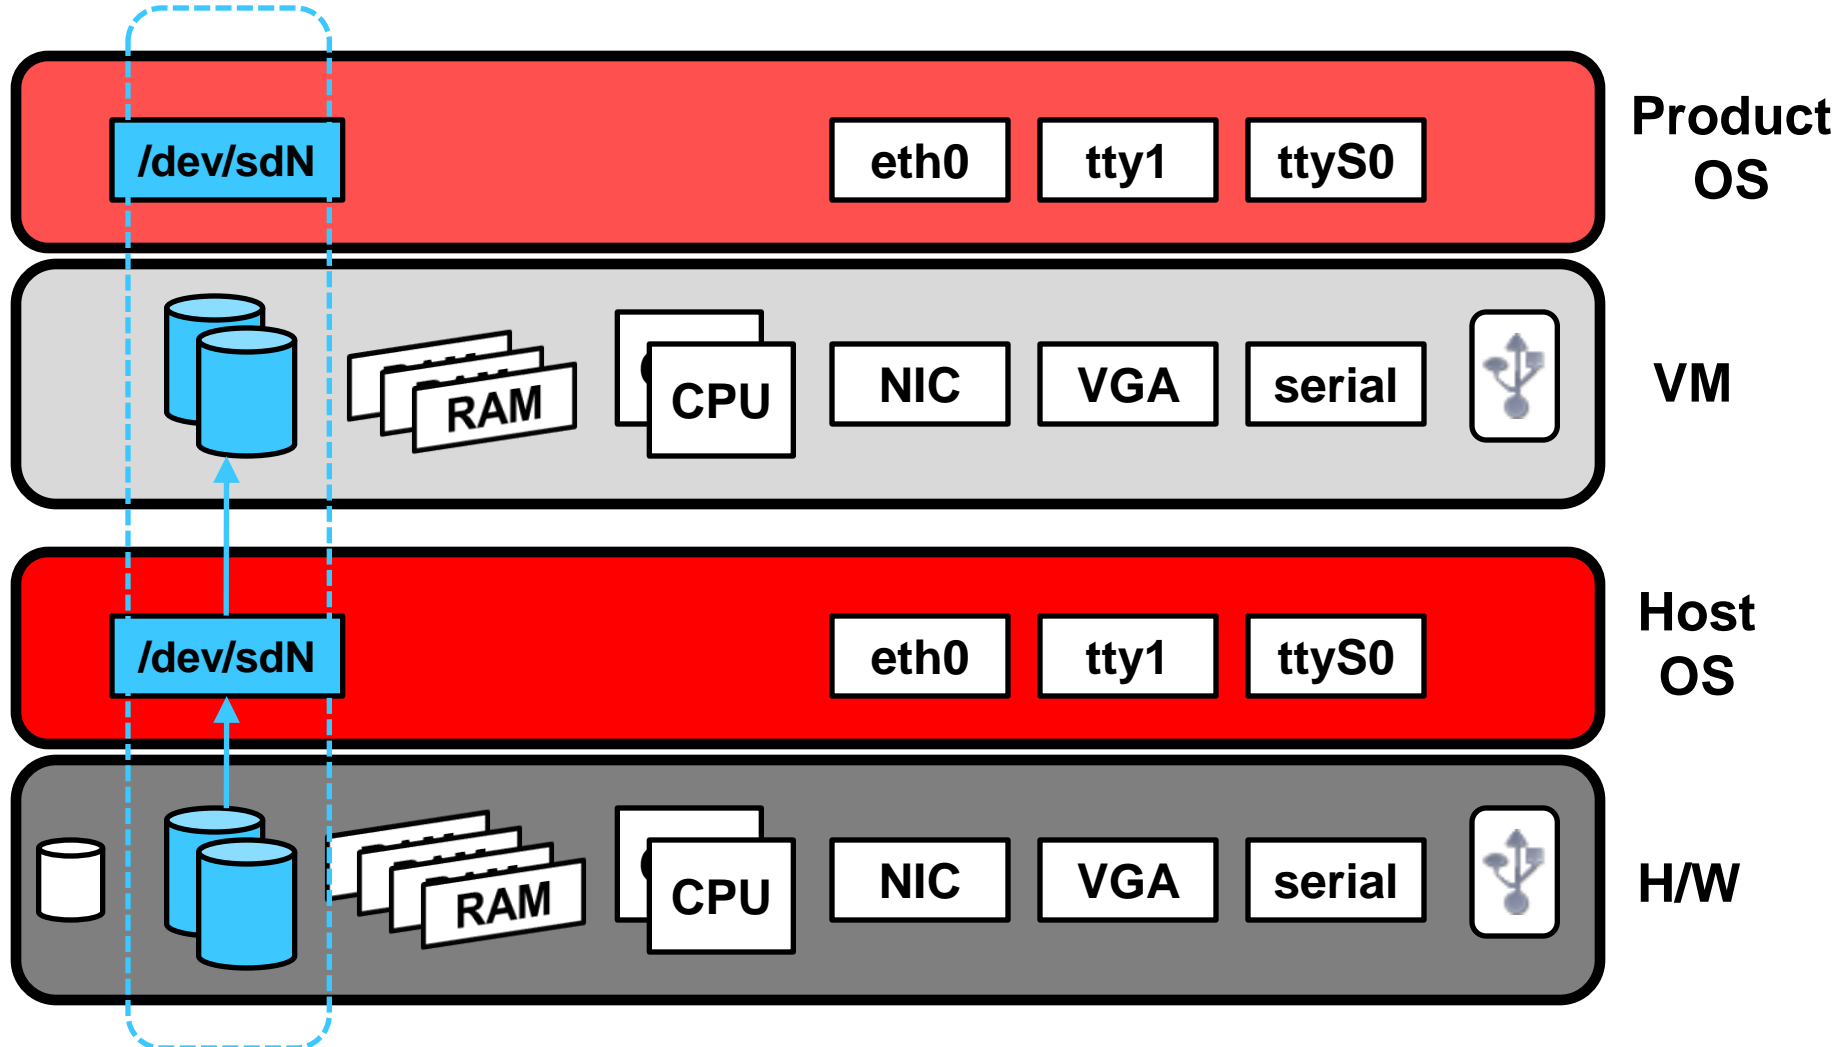
Storage Device for Host OS

- Internal USB key for host OS



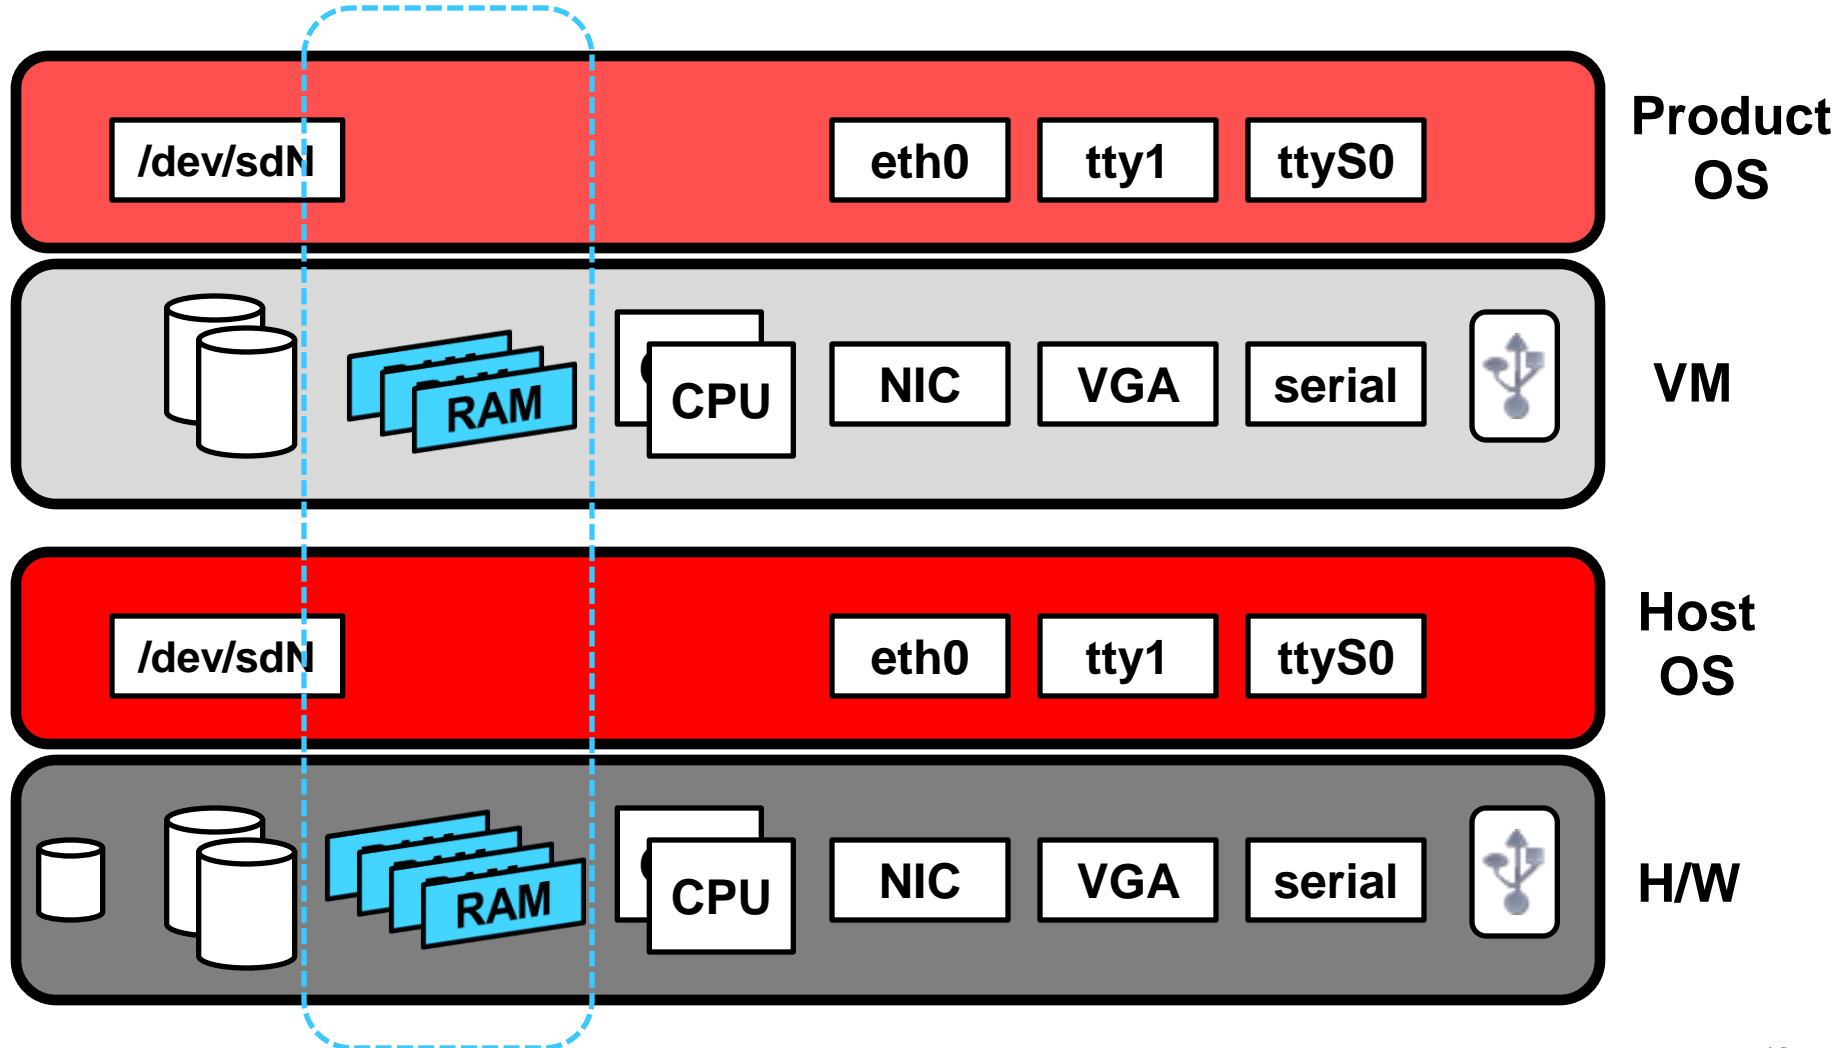
Hard Drives Given to VM

- Block devices in Host OS assigned as disks to VM



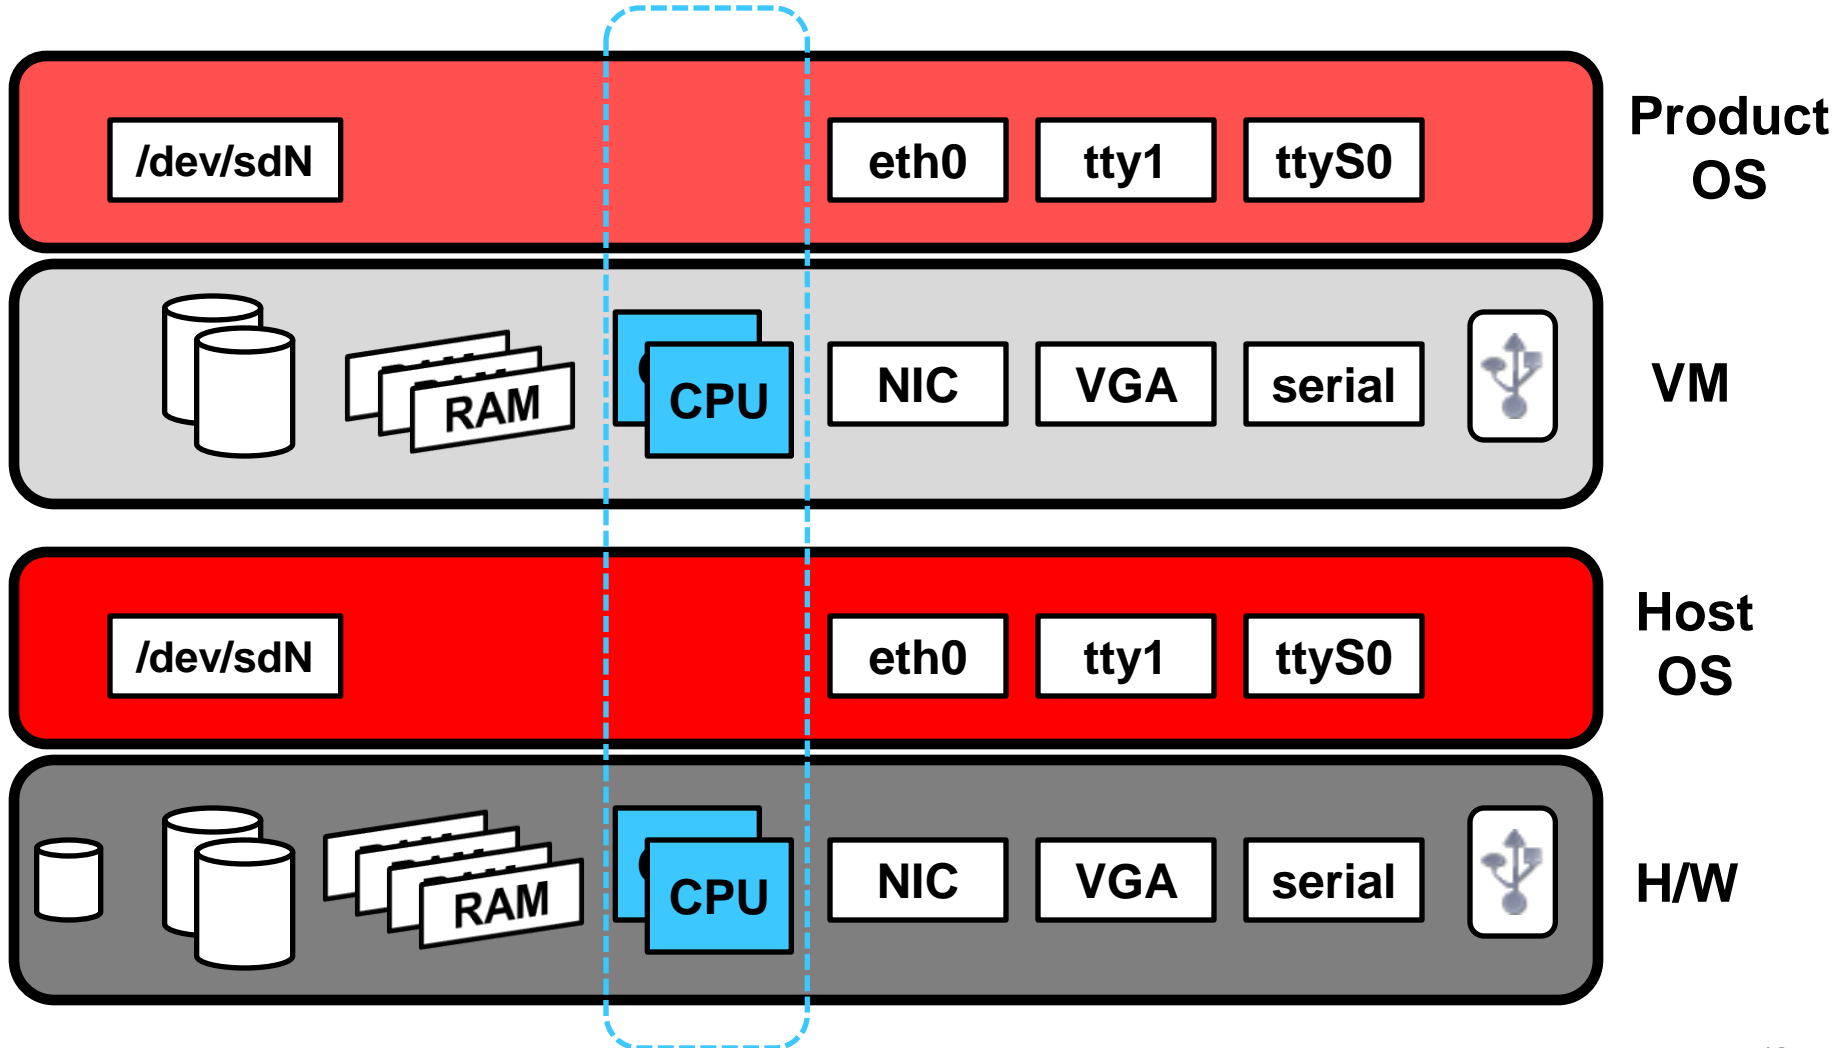
Memory Allocation

- Some RAM held back for Host OS



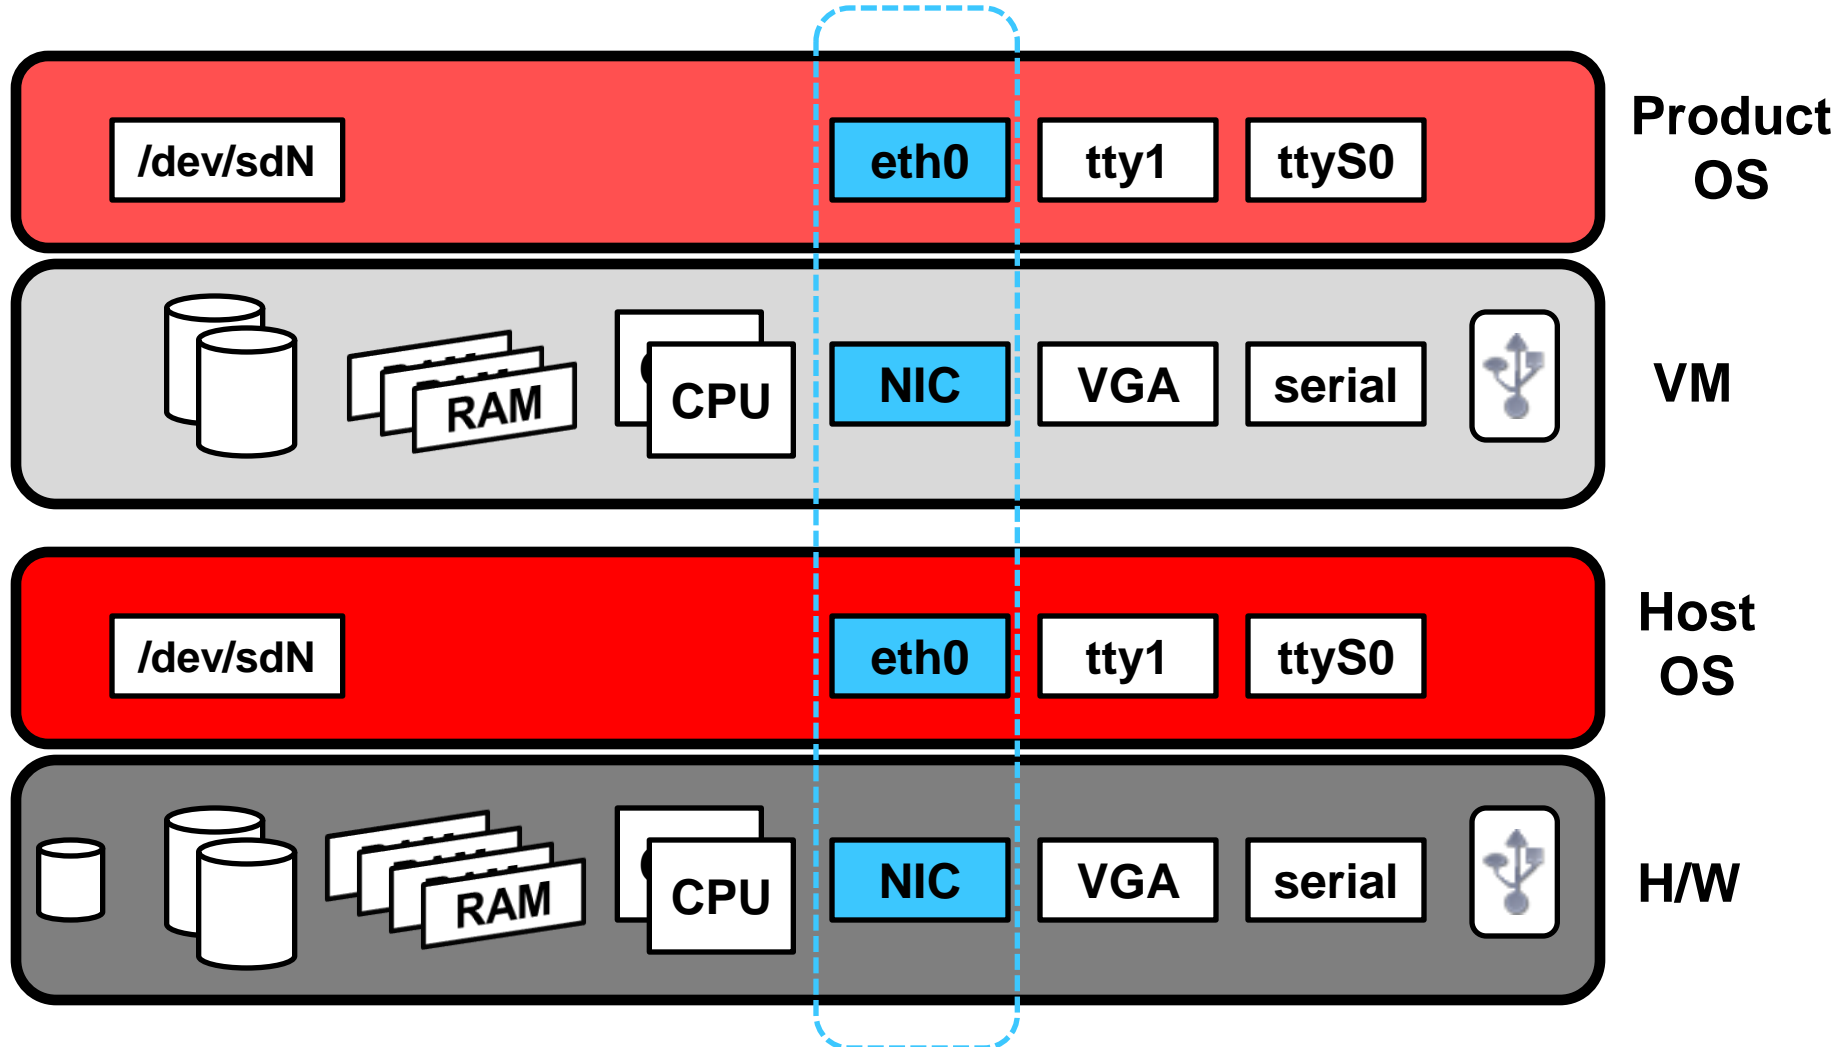
Processors

- # of Vcpus == # of Pcpus (excluding hyperthreads)



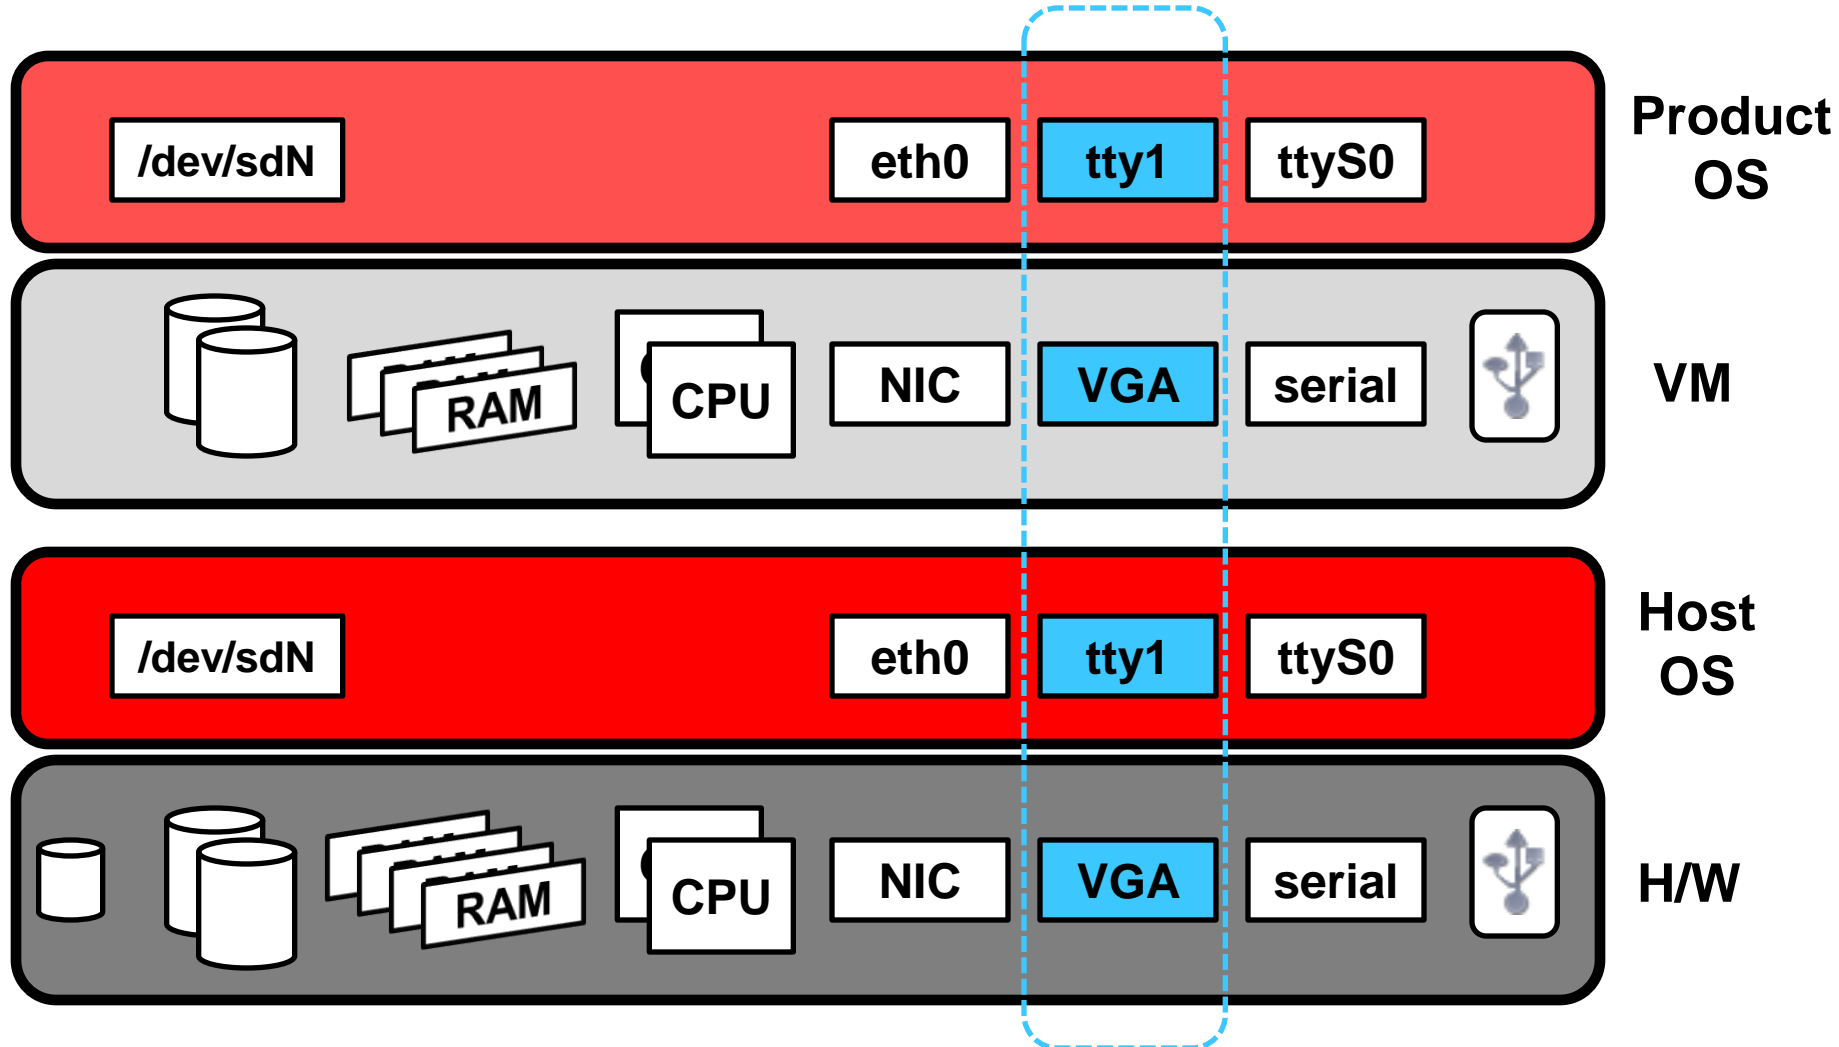
Networking

- First NIC of VM bridged to eth0 of host



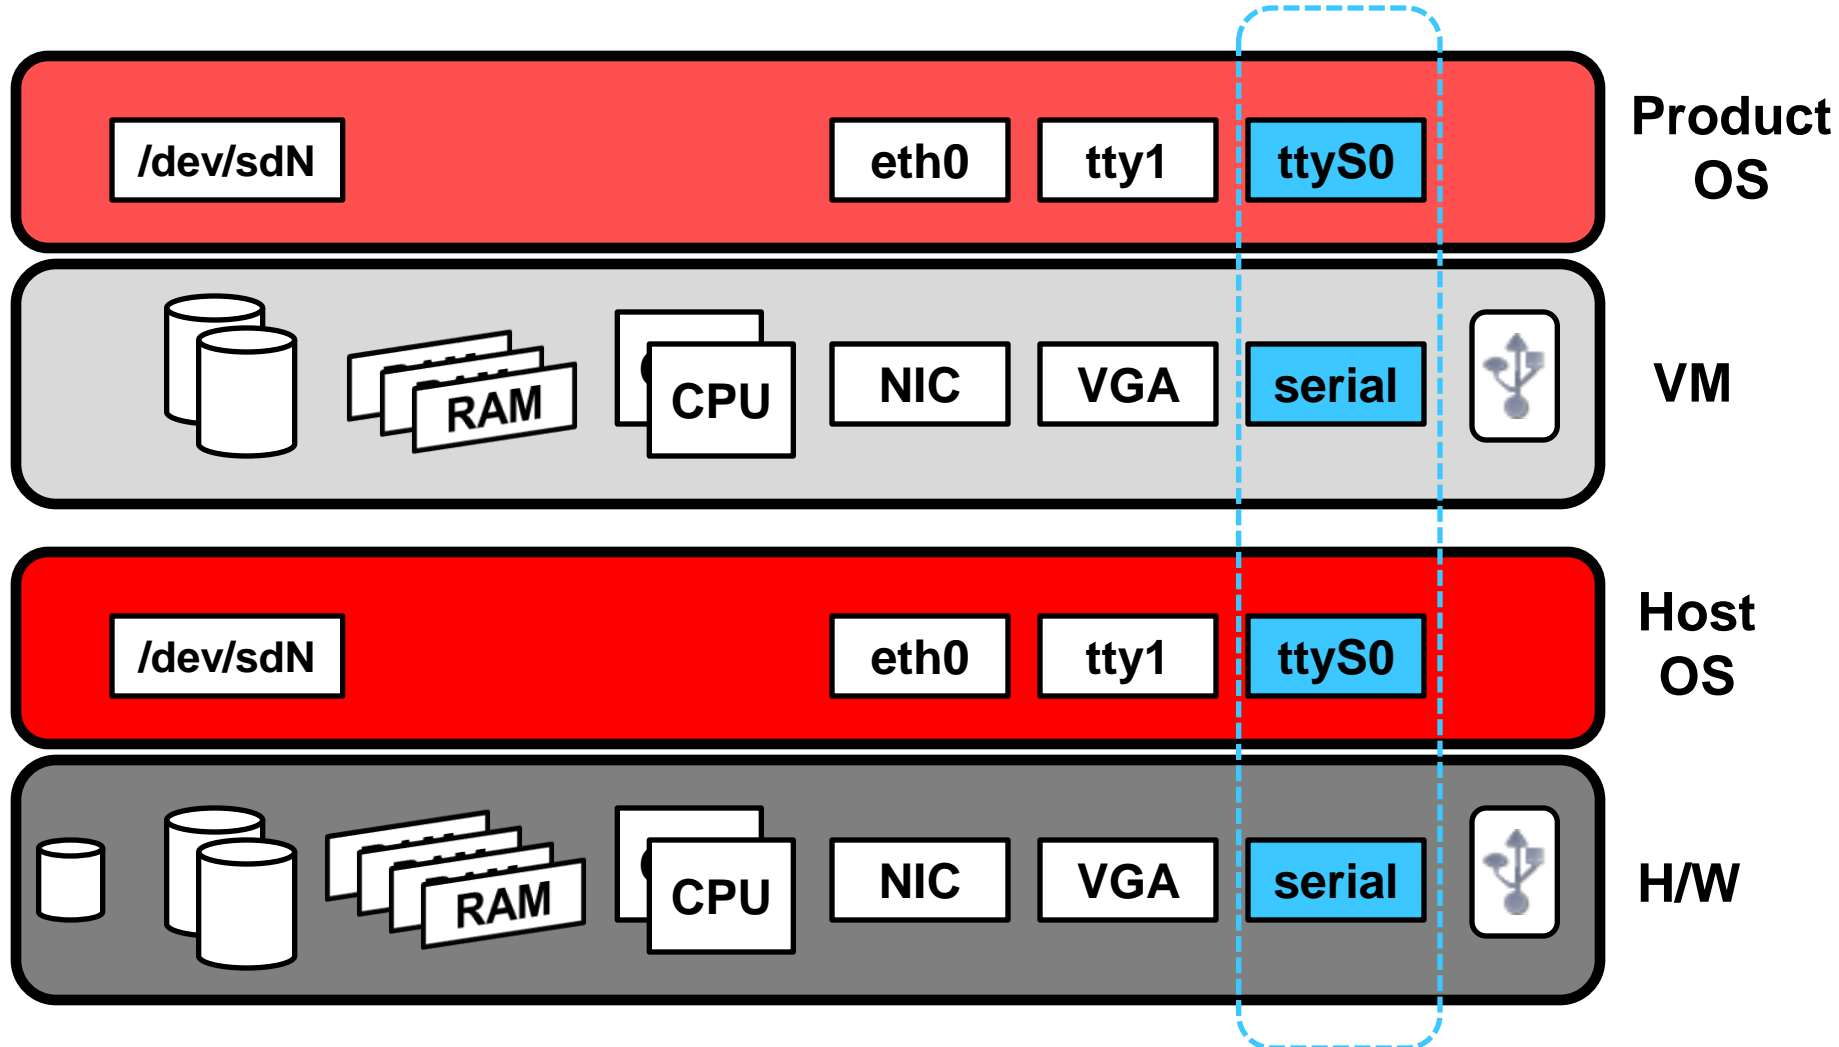
Console Display

- Console of VM connected to tty1 of host



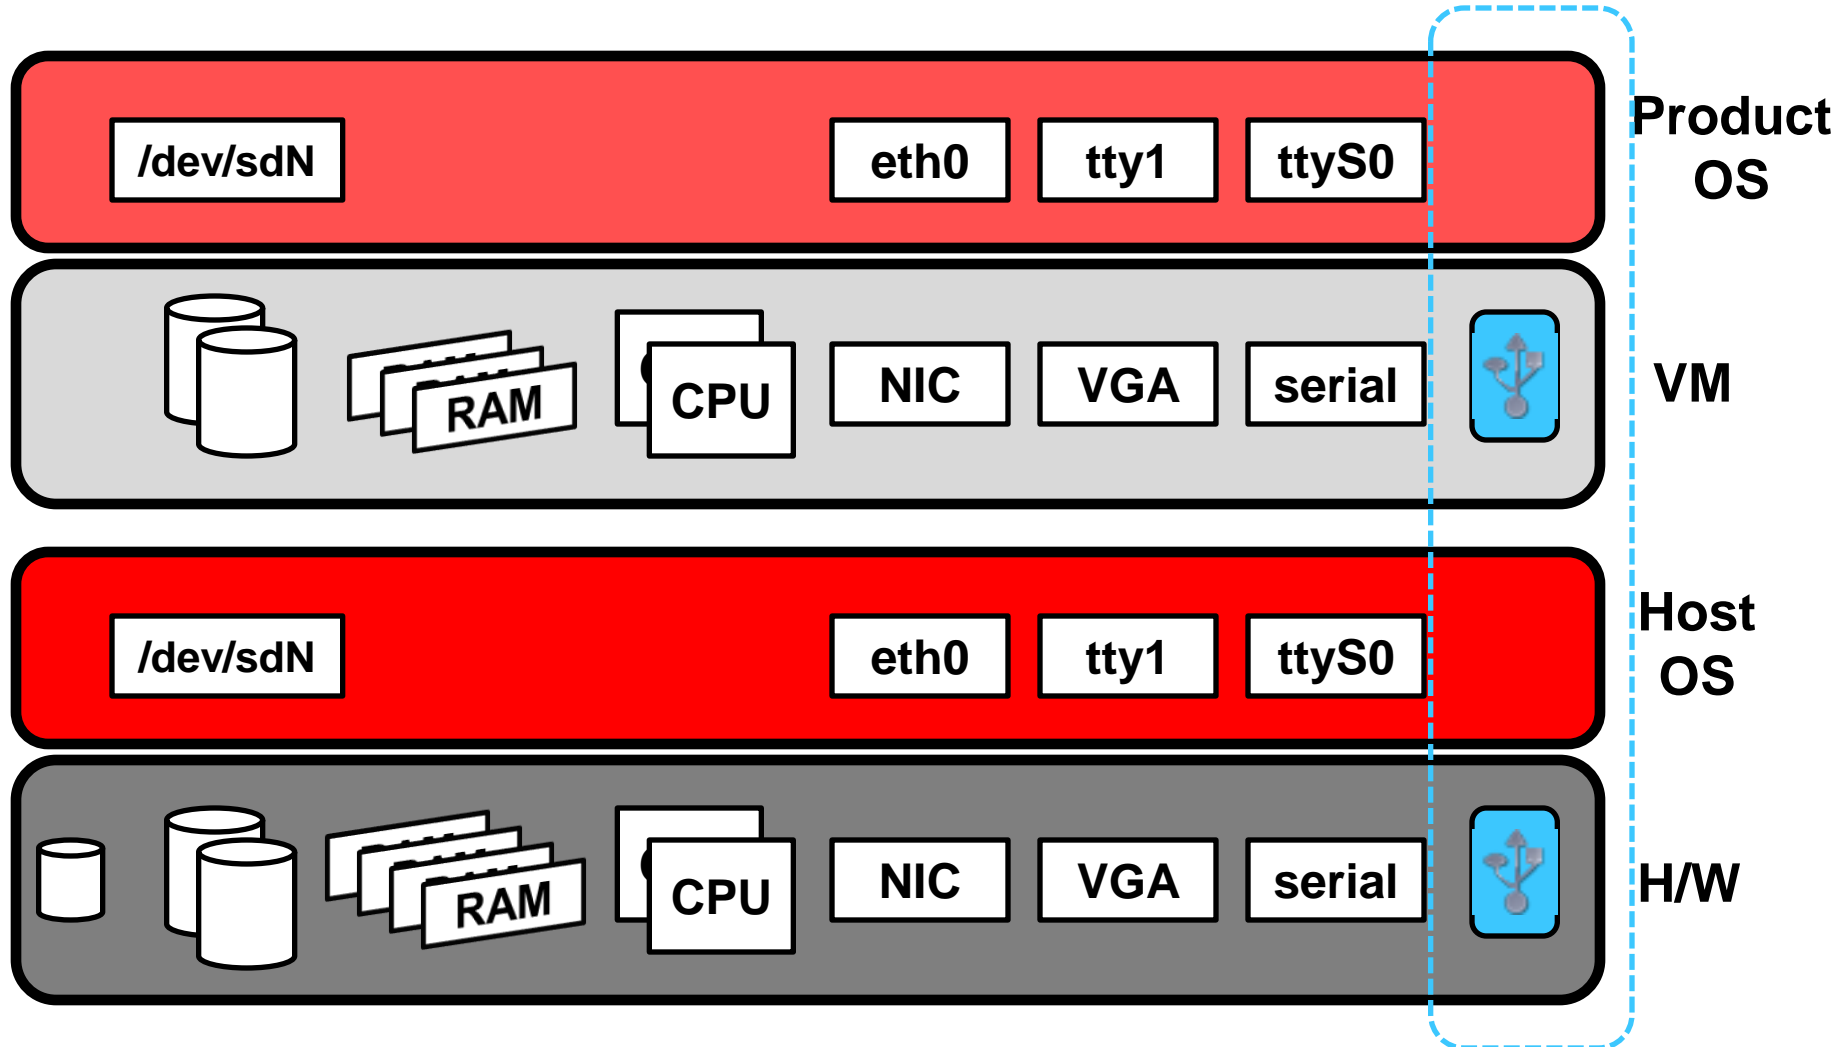
Serial Port

- Serial port of VM connected to HW



USB Devices

- USB devices passed to VM – a few exceptions



Marketing Constraints

- No change to end user experience
 - Installation, administration, monitoring, access
 - i.e., make HAL completely transparent to customer
- One DVD for product release that works for all supported hardware in the release
- HAL only installed on servers that need it
 - No user intervention required
- HAL removed upon upgrade to product release that recognizes H/W natively
 - No user intervention required
- Toggle between product releases enables/disables HAL
 - No user intervention required

Where Does KVM Fit In?

- KVM is the enabling technology for the HAL
 - behind the scenes
- qemu-kvm provides the virtual machine and device models
- KVM provides the virtualization efficiencies

Why KVM?

- KVM's architecture ideal for “embedded” use cases
 - Allows use of standard linux distribution as the host OS
 - Use of virtualization not relevant and can be hidden from end-user
- Host OS has same ‘look-and-feel’ as product OS
 - Install and runtime (development perspective) and diagnostics (customer support perspective)
 - Essentially creating a linux-on-linux stack – load kernel module, start userspace command
- Qemu and Linux are both powerful Swiss Army knives
 - Plenty of options to meet transparency requirements and maintain appliance-like design

Installation, device handling, console, VM management, etc

HAL Design Details



Design Details

- **Installation**
- Networking
- Product CLI
- SNMP Monitoring
- Syslog and Alarm Generation
- Devices – DVD and USB
- File Sharing
- Upgrades and Disabling HAL
- ...

Product Installation

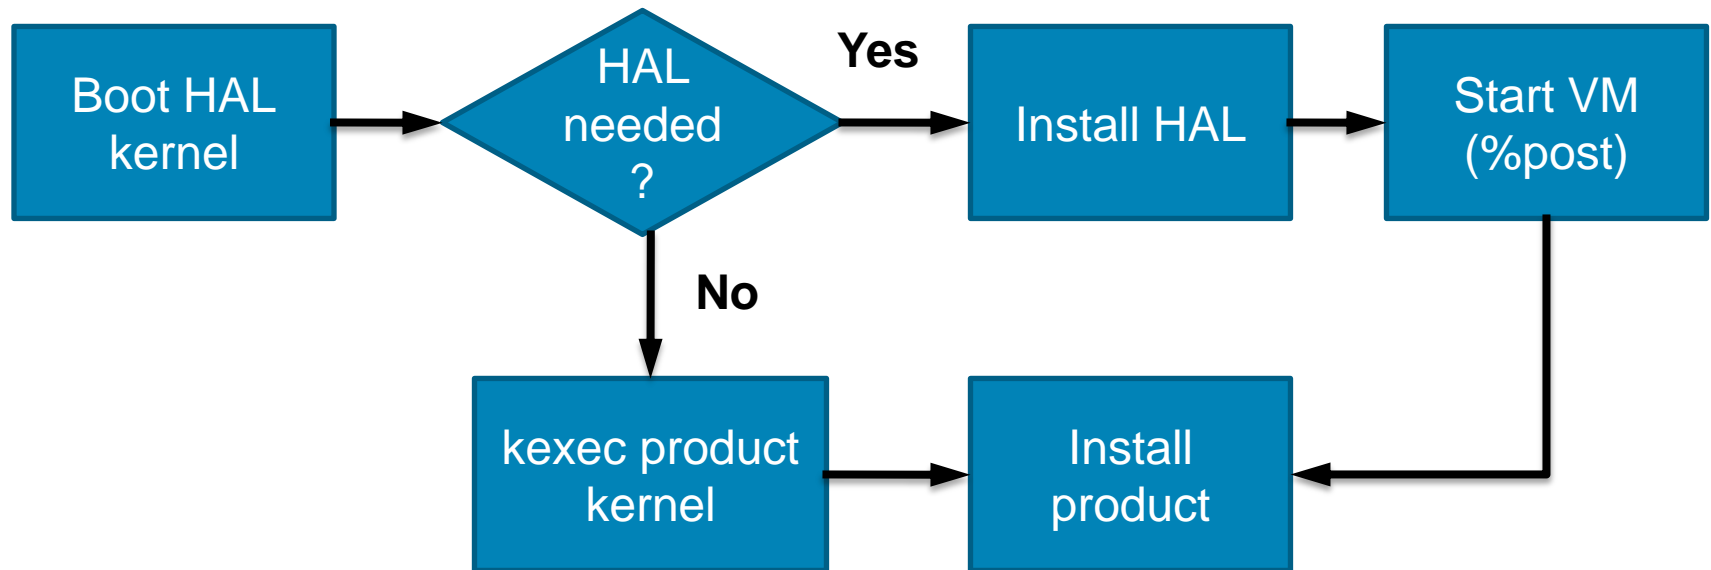
- Typical Install Sequence
 - Verify H/W is a supported platform
 - Update BIOS / firmware as needed
 - Configure BIOS and RAID
 - Install OS
 - Install product
 - Configure

Installation Impacts

- H/W aspects done during HAL install
 - Commands for H/W recognition, firmware updates, configuring RAID/BIOS, etc need to be moved to HAL
 - Updated for compatibility with Host OS
- HAL and product are two different OS installs
 - Each needs to be done within its context
 - Separate kernel + initrd.img for each OS
 - Co-exist on single DVD

Installation: One DVD For All Servers


- DVD boots HAL kernel + initrd
- Inspects H/W and decides if HAL is needed



- Non-HAL servers
 - Jump to product's OS happens quickly; install proceeds

Trampoline Kernel for DVD

- Some cases need a “trampoline” to jump between kernel versions
 - e.g., transitioning from RHEL5, 64-bit for HAL to RHEL3, 32-bit for product
- Trampoline is a very small kernel + initrd
 - No ACPI, no kernel modules
 - Just enough runtime env for basic platform detection
 - Takes less than 1 second to run and jump to next kernel



Installation with HAL

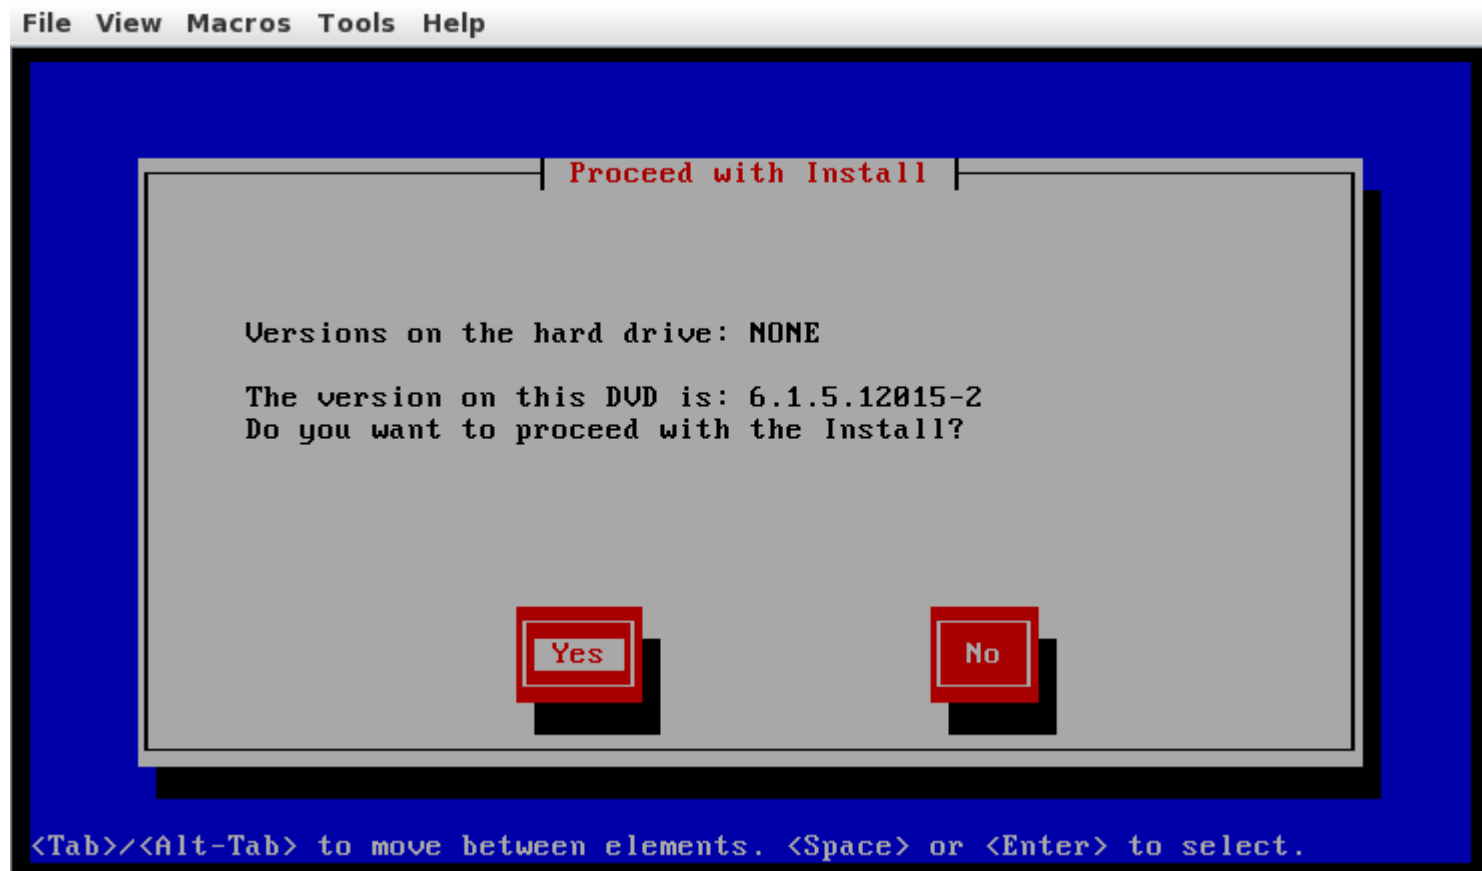
- H/W interaction done by HAL installer
 - Valid platform detection
 - BIOS, firmware updates
 - RAID/BIOS configuration
- Finds block device for Host OS
 - e.g., internal USB key
- Installs Host OS
- Starts VM for product install in post-install phase
 - Leverage `–kernel` and `–initrd` options of `qemu`
 - No reboot between install of HAL and product
 - Allows DVD to be used for install in VM without user intervention (ie., re-inserting)

Product Install

- Product install starts seamlessly after HAL install
 - VM console on host's tty1
 - qemu started with `–curses` display option
 - stdin/stdout set to `/dev/tty1`
- Same installation sequence as bare metal
 - Minus the H/W updates/configuration
- Maintains current user experience with console based installs
- Put installer into 'text' mode (e.g., text option for Anaconda)

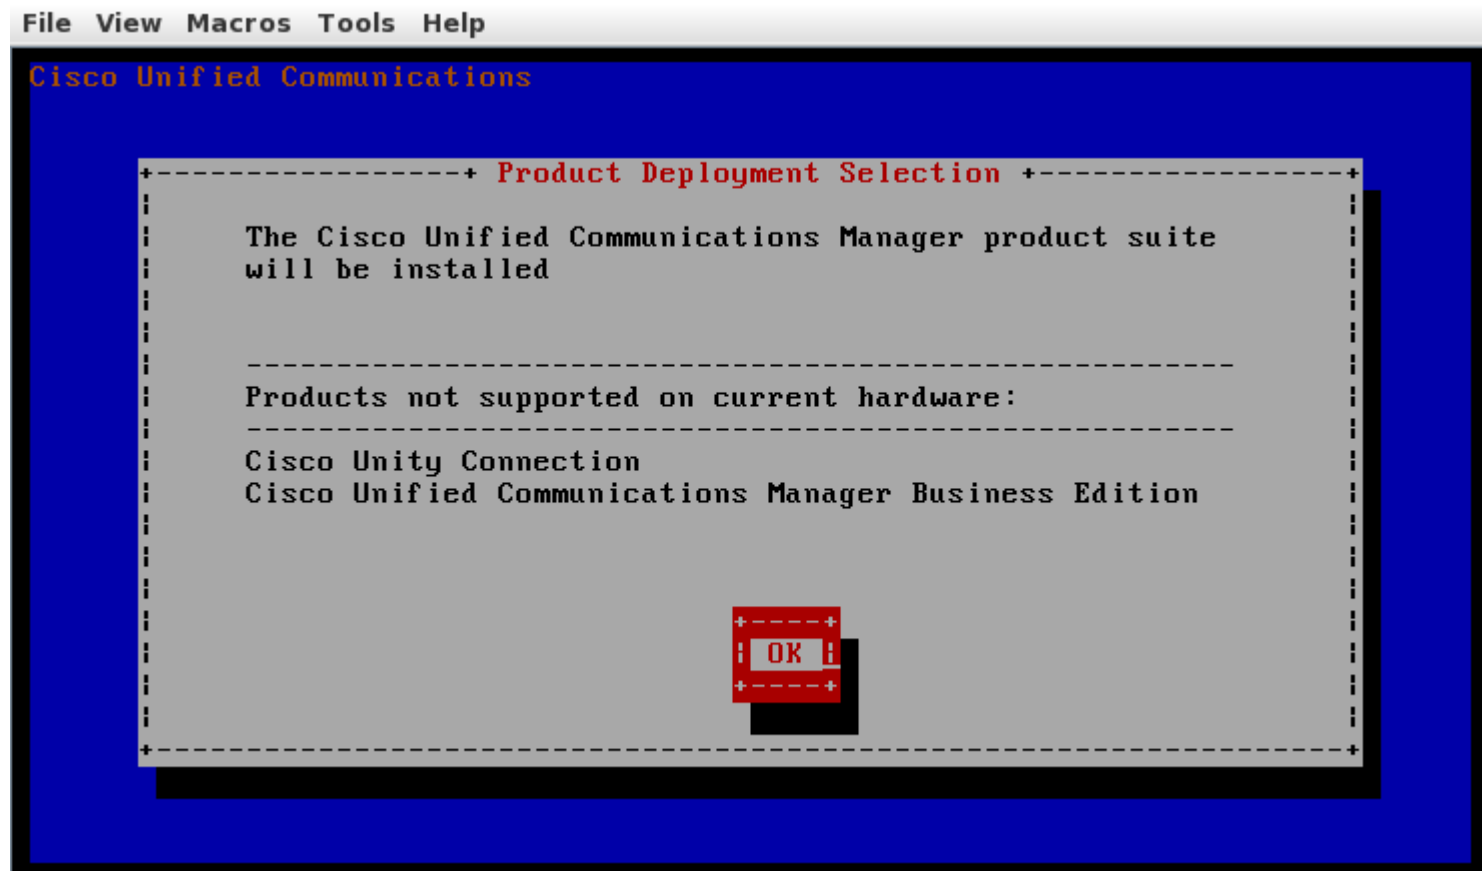
Initial Install Screen – HAL Installer

- Same sequence of input screens as bare metal
- First screen presented by HAL installer



Initial Product Install Screen

- Second input screen to user is from product installer
 - Slight loss in aesthetics due to qemu + curses path

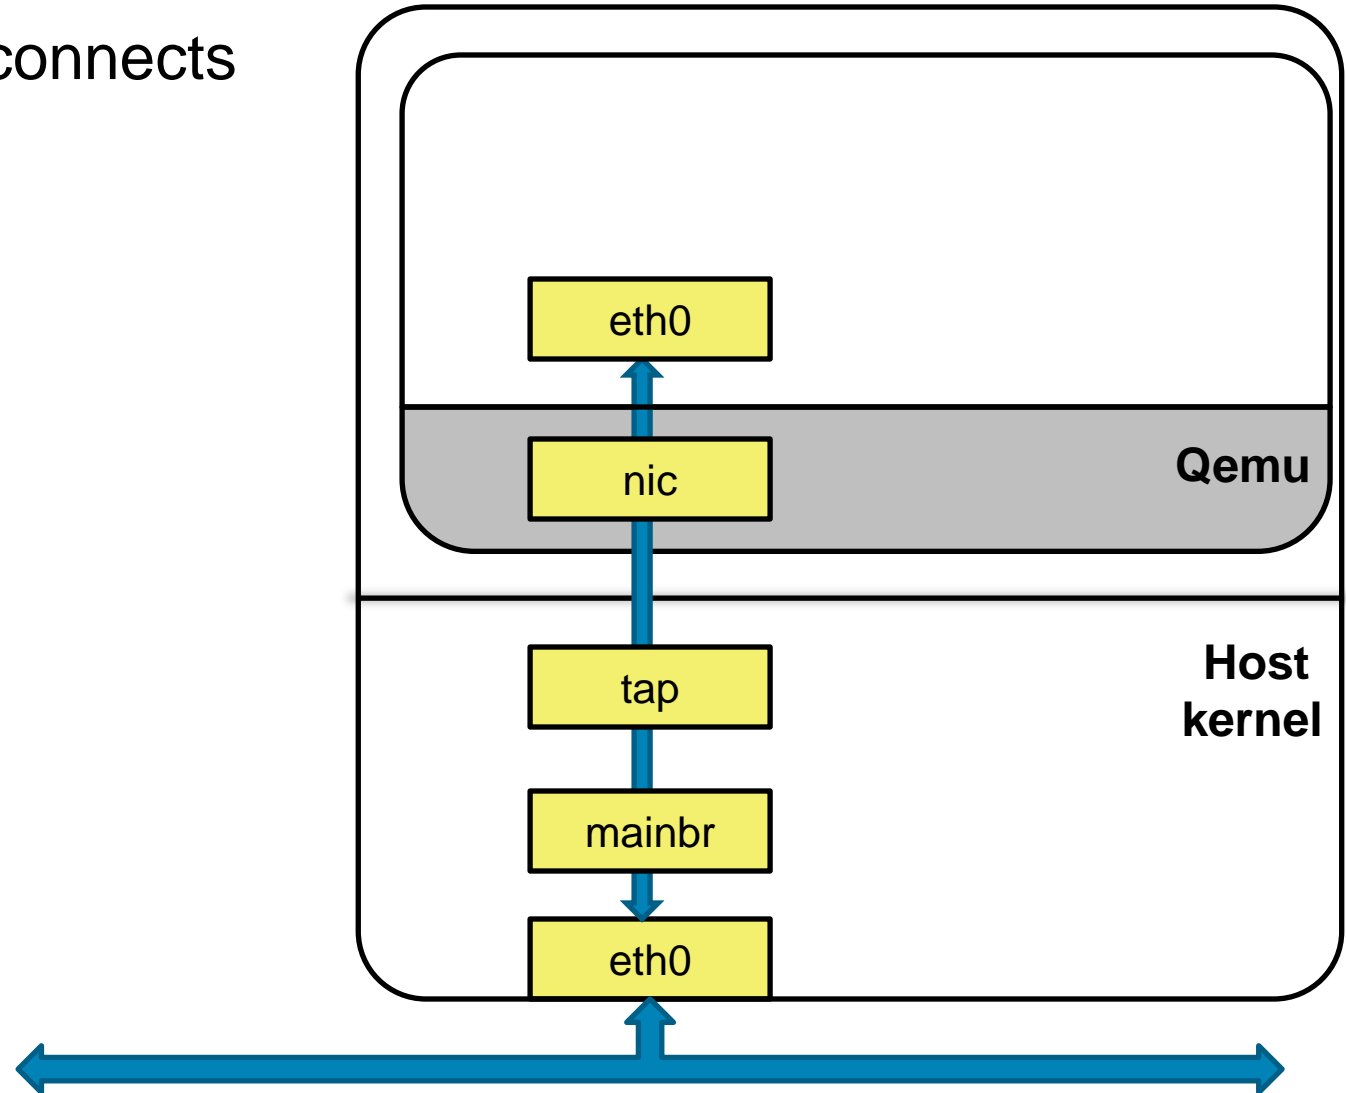


Design Details

- Installation
- **Networking**
- Product CLI
- SNMP Monitoring
- Syslog and Alarm Generation
- Devices – DVD and USB
- File Sharing
- Upgrades and Disabling HAL
- ...

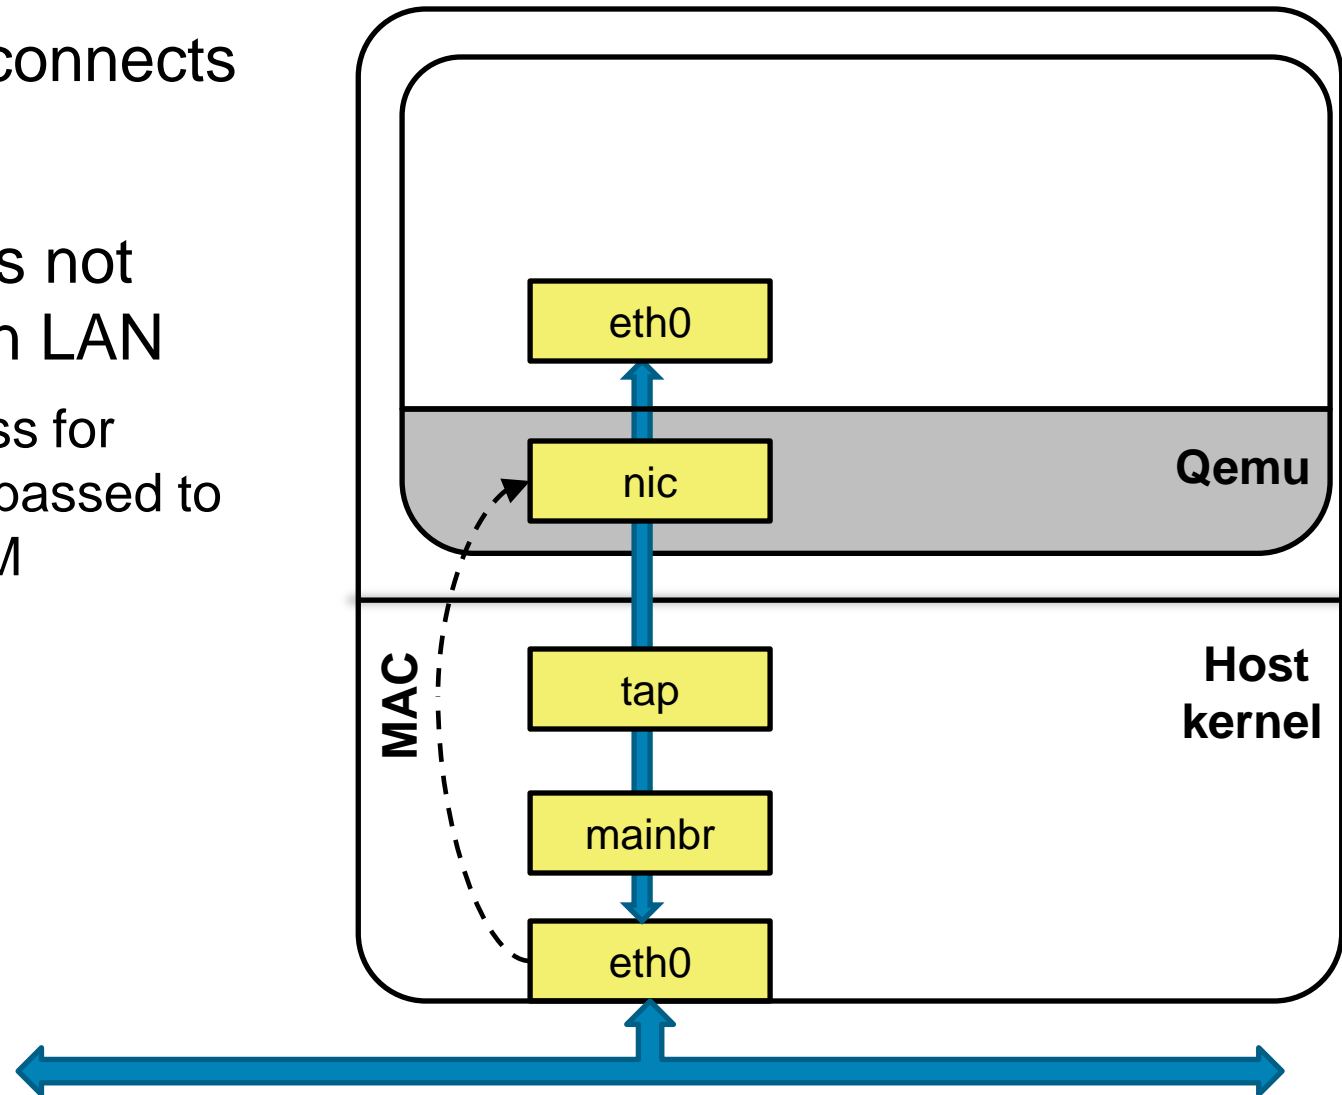
Networking

- Main bridge connects VM to LAN



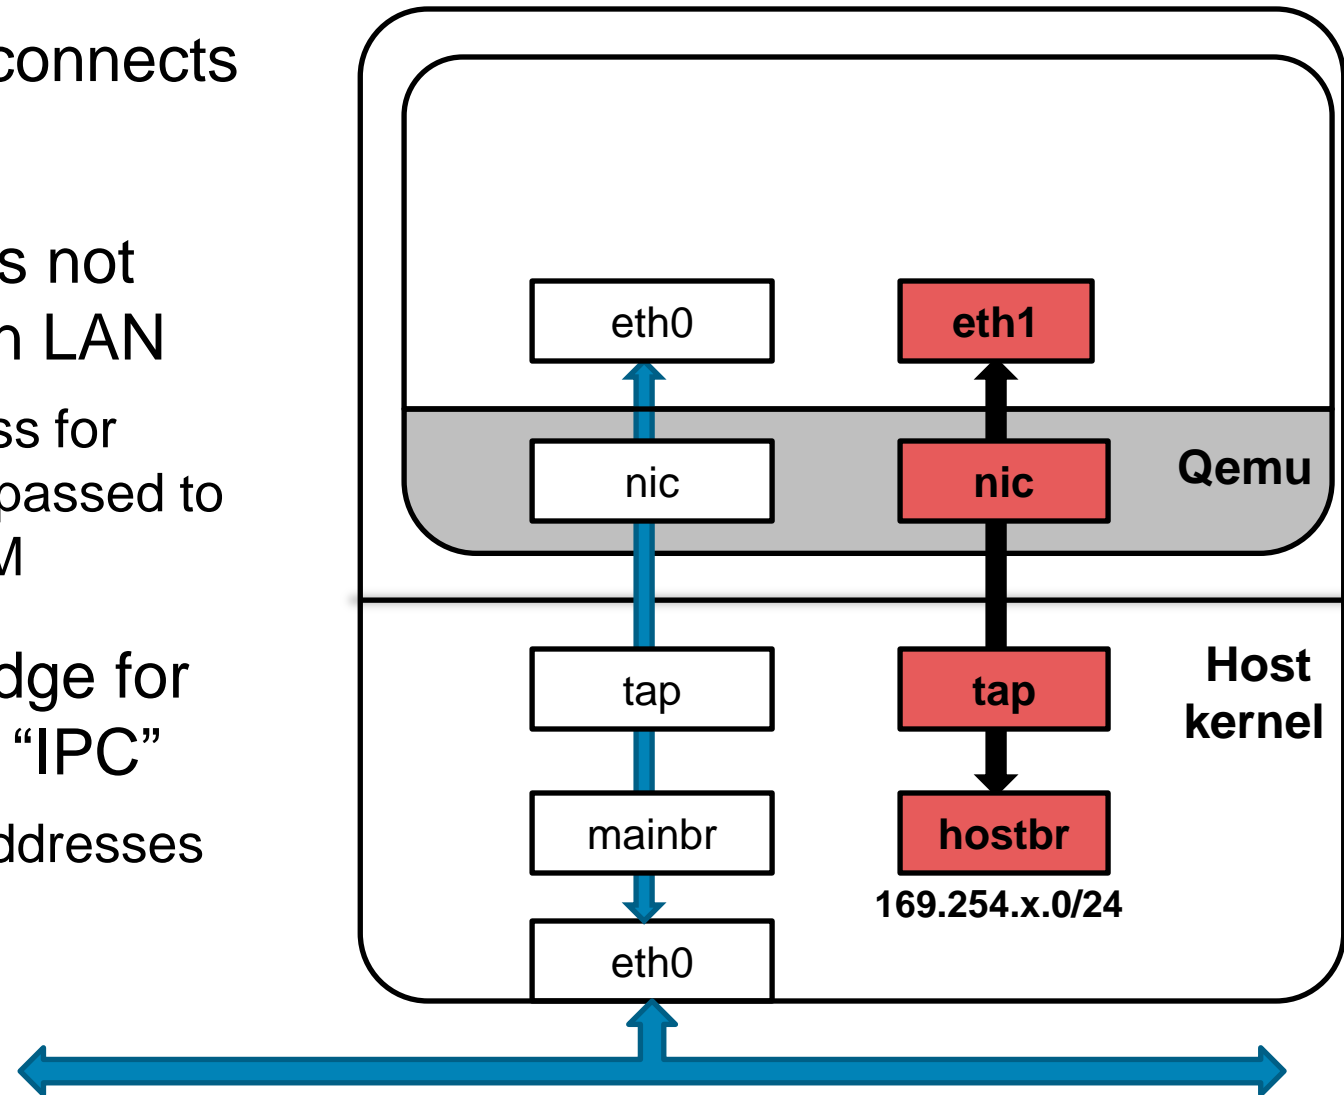
Networking

- Main bridge connects VM to LAN
- Host OS does not have an IP on LAN
 - MAC address for physical NIC passed to first NIC of VM



Networking

- Main bridge connects VM to LAN
- Host OS does not have an IP on LAN
 - MAC address for physical NIC passed to first NIC of VM
- Host-only bridge for Product-HAL “IPC”
 - Link-local addresses



Internal Network

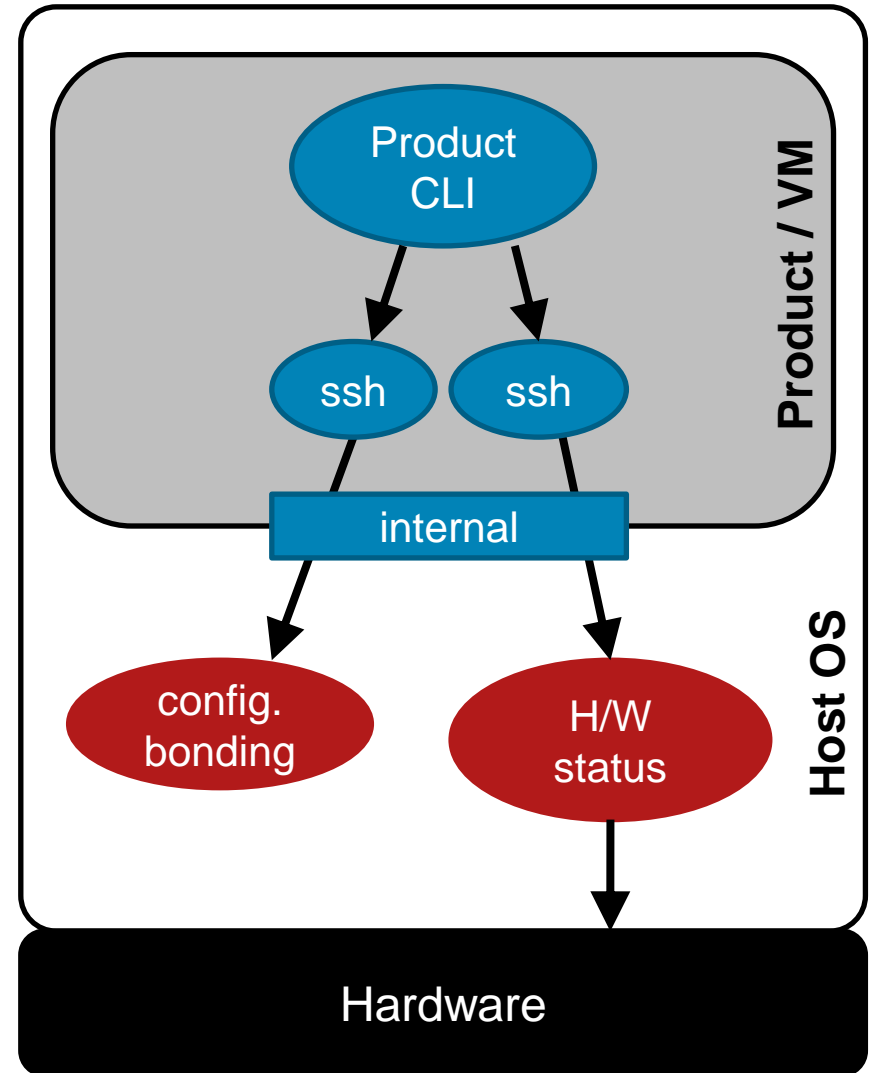
- Basis for HAL Transparency
- Allows access to H/W and HAL from within product
 - e.g., configuration & diagnostics
 - Standard protocols used for HAL-Product IPC: SNMP, syslog, ssh, NFS
- HAL-VM communications restricted to the “internal” network
 - Firewall rules in both layers restrict traffic for the link-local addresses to expected interfaces

Details

- Installation
- Networking
- **Product CLI**
- SNMP Monitoring
- Syslog and Alarm Generation
- Devices – DVD and USB
- File Sharing
- Upgrades and Disabling HAL
- ...

Product CLI

- Commands that configure or query H/W need to be run in Host OS
 - e.g., bonding, NIC or RAID status
- ssh over internal network
 - public key authentication
- CLI software modified to prepend HAL wrapper to backend commands
- Backend commands updated for Host OS

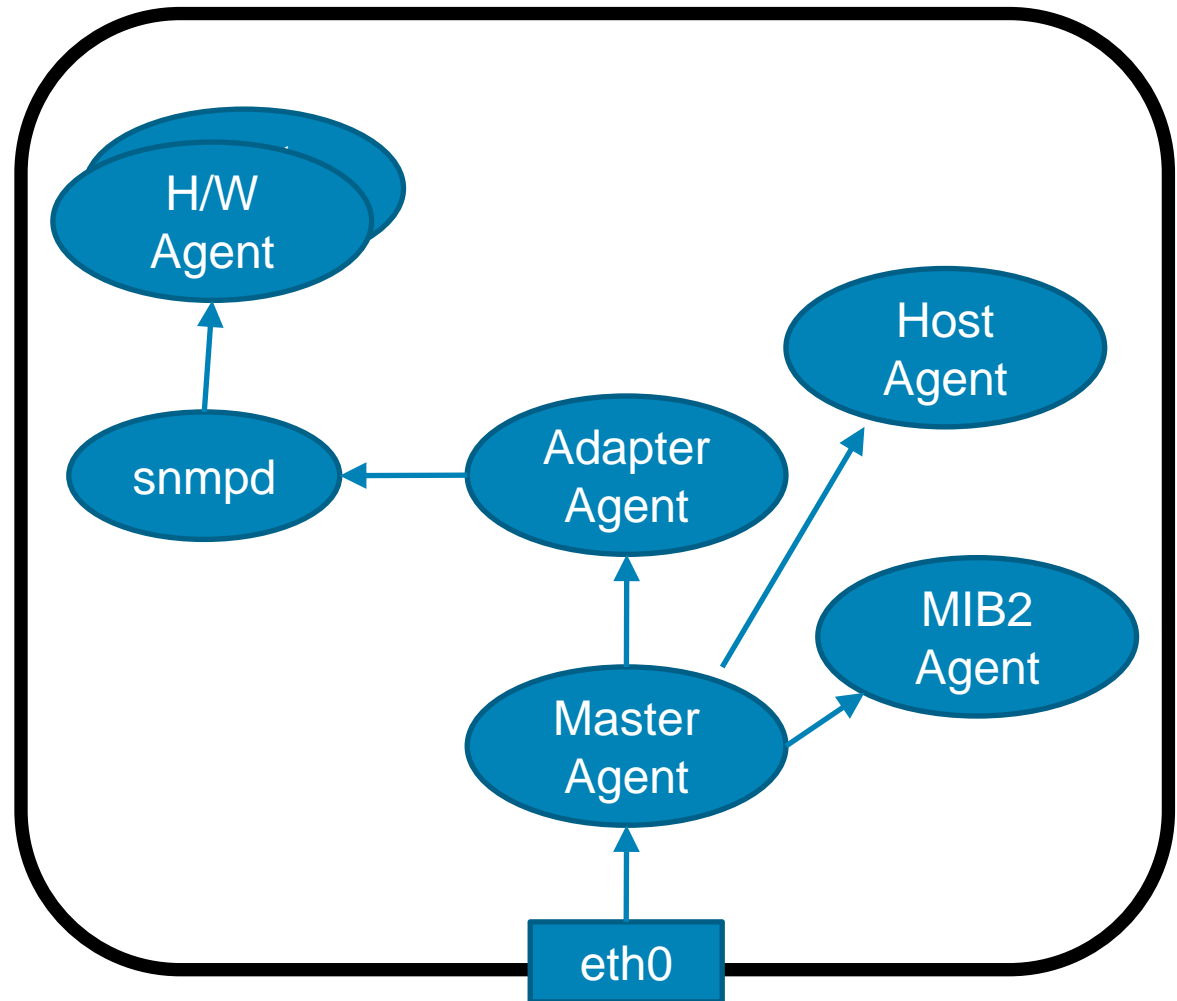


Design Details

- Installation
- Networking
- Product CLI
- **SNMP Monitoring**
- Syslog and Alarm Generation
- Devices – DVD and USB
- File Sharing
- Upgrades and Disabling HAL
- ...

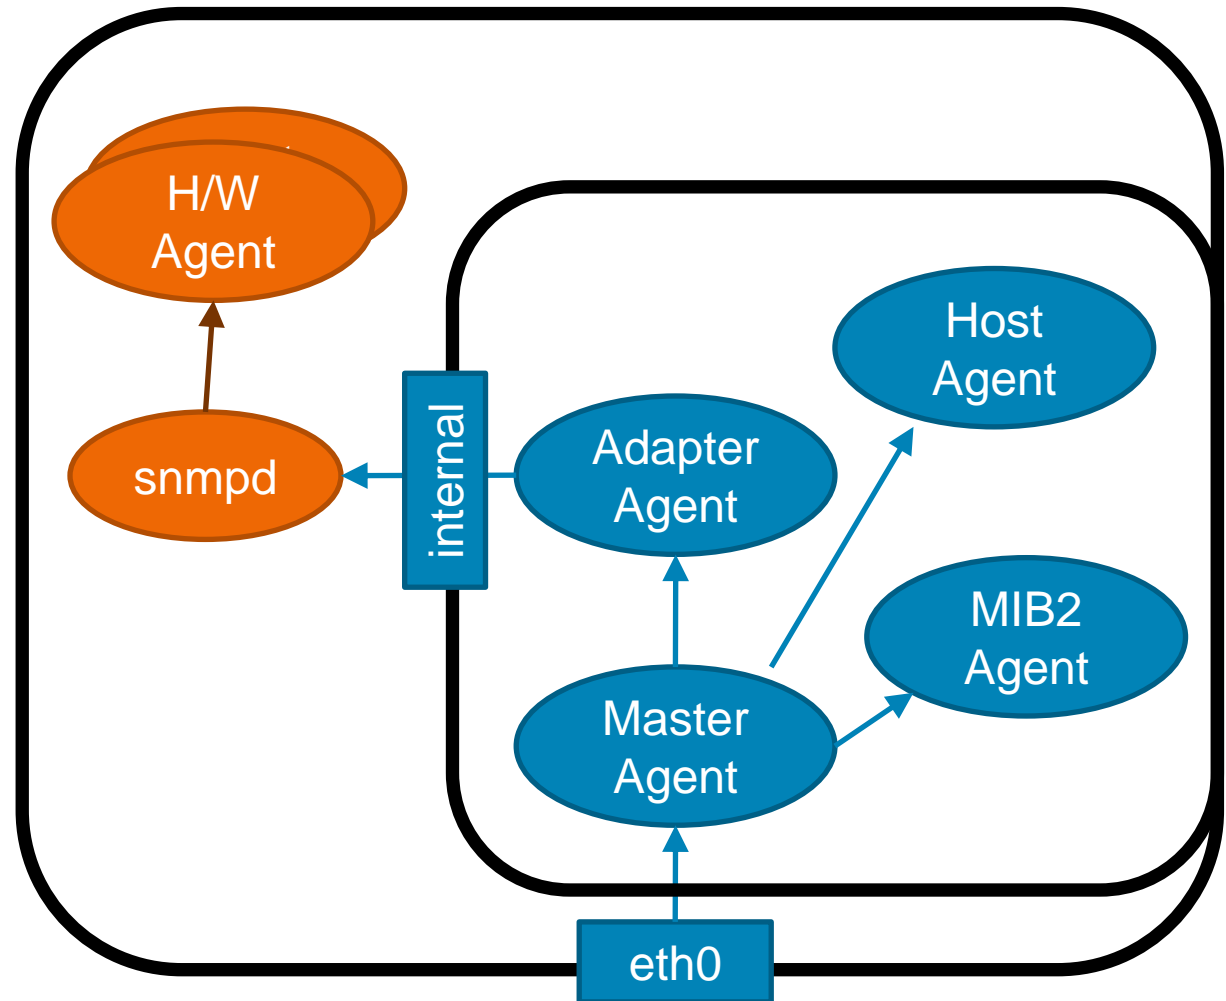
SNMP Architecture – no HAL

- Master Agent receives request
- Adapter agent used to connect to H/W agents from vendors
 - Forward done over loopback interface



SNMP Architecture – with HAL

- H/W-based agents need to run in the Host OS
- Configure adapter agent to forward over internal address

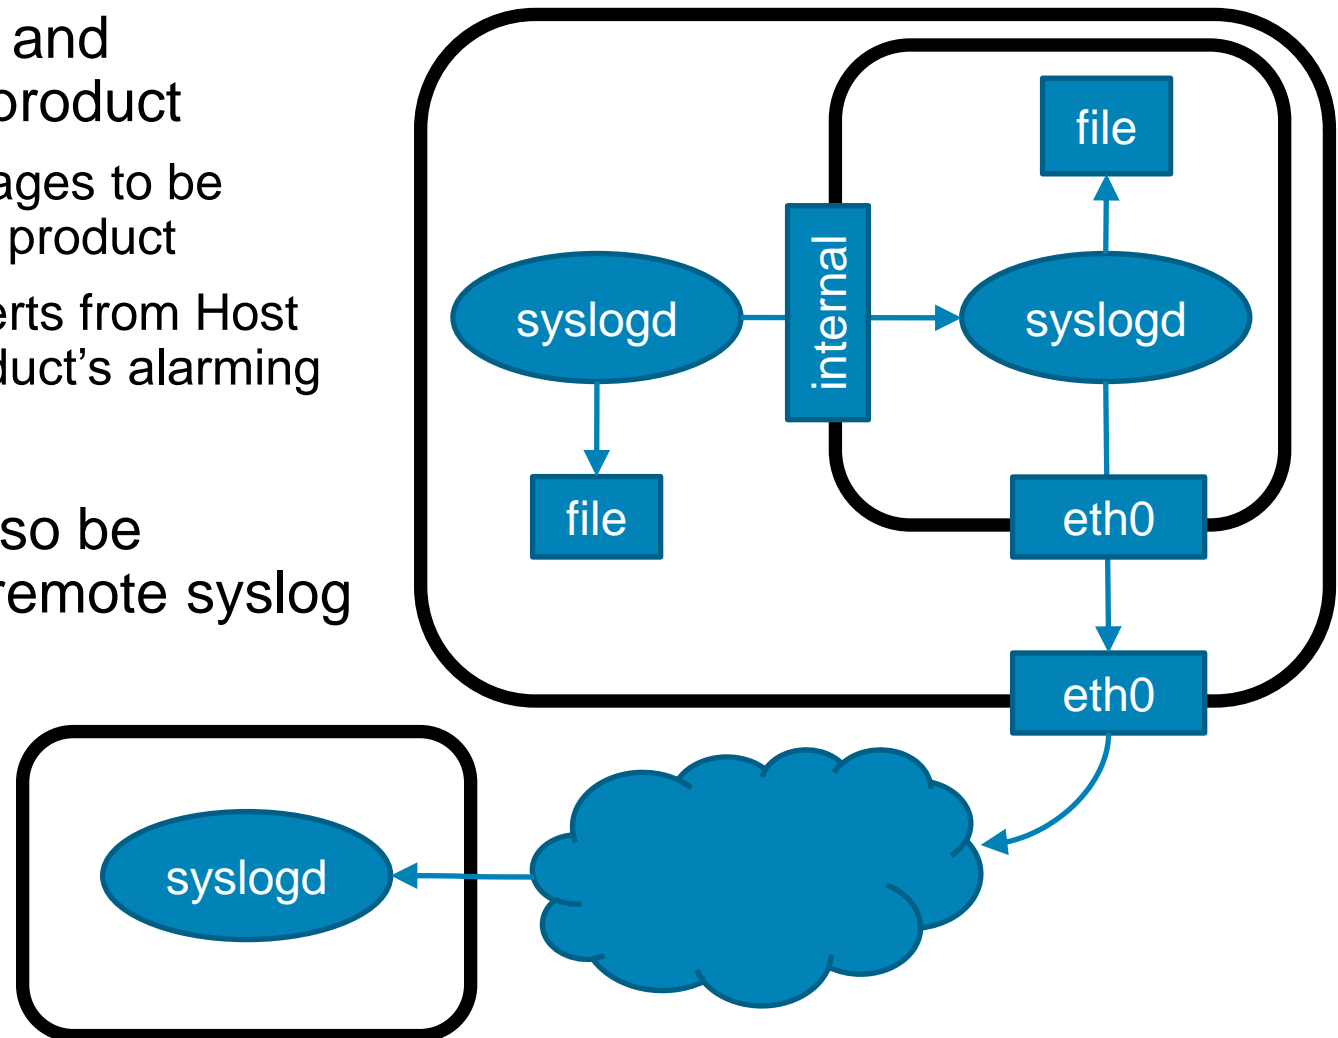


Design Details

- Installation
- Networking
- Product CLI
- SNMP Monitoring
- **Syslog and Alarm Generation**
- Devices – DVD and USB
- File Sharing
- Upgrades and Disabling HAL
- ...

Syslog Forwarding

- syslog entries from HAL written locally and forwarded to product
 - Allow messages to be processed by product
 - e.g., HW alerts from Host OS enter product's alarming infrastructure
- Entries can also be forwarded to remote syslog daemon



Design Details

- Installation
- Networking
- Product CLI
- SNMP Monitoring
- Syslog and Alarm Generation
- **Devices – DVD and USB**
- File Sharing
- Upgrades and Disabling HAL
- ...

Device “Passthrough”

- Emulation sense, not VT-d
- DVD device opened by VM when it boots
 - access to DVD for upgrades or other functions (e.g., password reset)
 - /dev/cdrom in Host OS maps to /dev/cdrom in VM
- H/W serial port connected to VM serial port via character device in Host OS
- USB devices added to and removed from VM when inserted/removed from hardware
 - udev script in Host OS
 - e.g., USB serial cable, audio device, storage device

Design Details

- Installation
- Networking
- Product CLI
- SNMP Monitoring
- Syslog and Alarm Generation
- Devices – DVD and USB
- **File Sharing**
- Upgrades and Disabling HAL
- ...

File Sharing

- Parts of HAL filesystem exported to VM using NFS
 - restricted to internal network only
- Allows access to HAL log files from within VM
- Allows product to push files to the HAL
 - HAL, BIOS, firmware upgrades

Design Details

- Installation
- Networking
- Product CLI
- SNMP Monitoring
- Syslog and Alarm Generation
- Devices – DVD and USB
- File Sharing
- **Upgrades and Disabling HAL**
- ...

Upgrades

- Dual boot partitions allows toggle between product releases
 - Use dual boot partitions for HAL too
- HAL can be upgraded as part of ongoing maintenance releases of product
 - From product perspective HAL update similar to BIOS or firmware update
- HAL can be disabled if new product version does not require it
 - Product upgrade to version with native support
 - e.g., Toggle boot order (for host OS running on USB key)

Automatic Enable / Disable of HAL

- HAL only installed on servers that need it
- Use Case: current product version requires HAL, upgrade to new version that recognizes H/W natively
 - HAL disabled on toggle to new version
 - On revert to older version, HAL re-enabled
- HAL runs from separate storage device
 - e.g., internal USB key
 - HAL enabled / disabled by modifying boot order
- Without HAL Product boots directly from harddrives
 - Recall: harddrives passed to VM, boot loader installed in typical fashion

Design Details

- Installation
- Networking
- Product CLI
- SNMP Monitoring
- Syslog and Alarm Generation
- Devices – DVD and USB
- File Sharing
- Upgrades and Disabling HAL
- ...

HAL Diagnostics

- HAL performance data collected using collectd
 - System resource usage (storage, cpu, RAM)
 - Per-process statistics for processes running in Host OS
- Data sent to collectd peer running in VM
 - Provides access to HAL performance, trending data within product
- Root-level account can be mirrored from VM to HAL
- Syslog messages forwarded to syslog in VM

Restart & Shutdown

- HAL follows VM for “commanded” actions by user
 - if VM is restarted (ie., rebooted) HAL also reboots
 - If VM is shut down HAL also shuts down
- VM lockups
 - Guest manager running in Host OS exchanges heartbeats every 60 seconds with counterpart running in VM
 - VM restarted after N missed heartbeats
- HAL lockups
 - Handled by H/W functionality (e.g., ASR for HP)

HAL Security

- Minimal footprint OS
- VM runs as non-root user
- Leverage SELinux in Host OS
- Root disabled
- No default passwords
 - ssh public-key authentication for automated commands
 - sync root-level accounts when created/deleted within product
- HAL upgrades/patches applied via product
 - Like BIOS/firmware update

Performance

- Nehalem class processors
 - Overhead of virtualization layer is relatively low
- Newer hardware compensates for virtualization overhead
- Can leverage standard Linux OS features as needed
 - e.g., cpu affinity, huge pages, ksm

Net Results

- Transparency
 - Not 100%, but darn close – e.g., Host OS boot messages
 - Primary goal met - No interaction or intervention required from end user
- A Few Data Inconsistencies – e.g., SNMP
 - MIBs from H/W vendors have some overlaps with standard MIBs
 - H/W agents and standard agents run in different OS'es → data retrieved from two different views of the hardware

Advantages of HAL Approach

- More cost-effective approach to meeting preferred product timelines
 - Lifetime buys – money-based solution
 - OS backports – staff-based solution
 - HAL – technology-based solution
- Smaller development and test effort than OS backport
 - Changes primarily limited to new H/W
- Maintains closed-box, appliance model
- Forward looking
 - Leverages newer OS for HAL
 - Early experience with next OS that product will eventually move to with code re-use
- Re-usable design
 - Once hooks exist within product can be leveraged in future releases

Final Thoughts

- Not intended to allow products stay on an outdated OS indefinitely
 - Other factors drive need to change OS
 - e.g., Support for 3rd party code, security errata
- Lot of details had to be omitted for brevity
 - Please ask if you have questions

