



# Building a Better NOS with Linux and switchdev

netdev 0x12 - July 12, 2018

---

David Ahern | Shrijeet Mukherjee



# Agenda

---

- What is this Whitebox / Disaggregation / Open Networking {r}evolution?
- Evolution of Network Operating Systems
  - Increasing use of Linux
- The Next Step
  - Linux as the OS of the data center with ASIC drivers in the kernel

# Legacy Networking

---



# Whitebox Switches and Disaggregation

---





# Open Networking

---

- Networking Operating Systems still silos
  - Can pick a vendor, but still highly dependent on vendor for service and consulting
- What does “Open” mean?
  - Shell access?
  - Able to run Linux commands at a shell prompt?
  - Able to run networking programs in the control plane?

# Open Networking with Linux




- Enable innovation
  - Solving network problems in ways unimagined by NOS or ASIC vendors
- Flexibility and reconfigurability are important – **essential** characteristics
  - Workloads, performance demands and characteristics can change quickly
  - Need a network that can adapt
- A big blob of highly interdependent processes is not flexible
  - No insights into each component means no flexibility
- Linux is about building blocks
  - Building blocks provide ultimate flexibility and reconfigurability



# Networking Solutions Constantly Evolving


---





# Switches Today Have Similarities to Servers / Hosts

- Hardware components
  - Commodity CPU (x86, arm, ppc)
  - Storage devices – USB, SSD
  - Management NIC
  - Multiple “data plane” ports
- Servicing
  - Logging in remotely, diagnosing problems with familiar tools
- And, Linux is the primary OS
- Operational Model
  - Configuring interfaces and services
  - Monitoring – e.g., interface statistics, events, ...

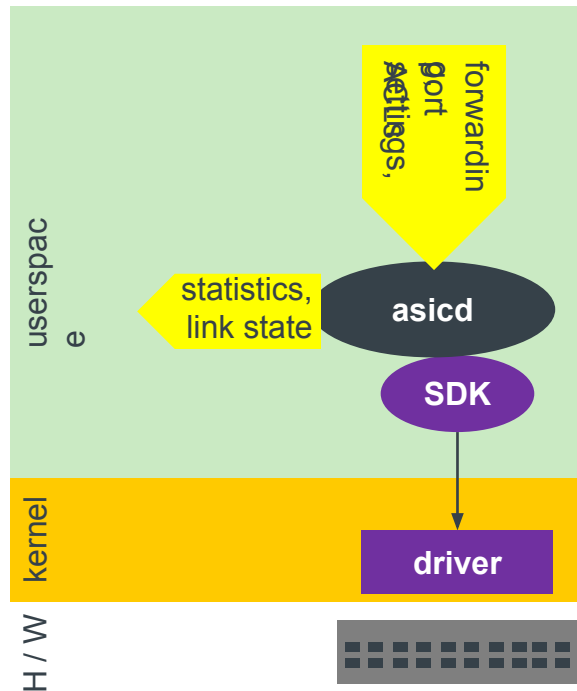




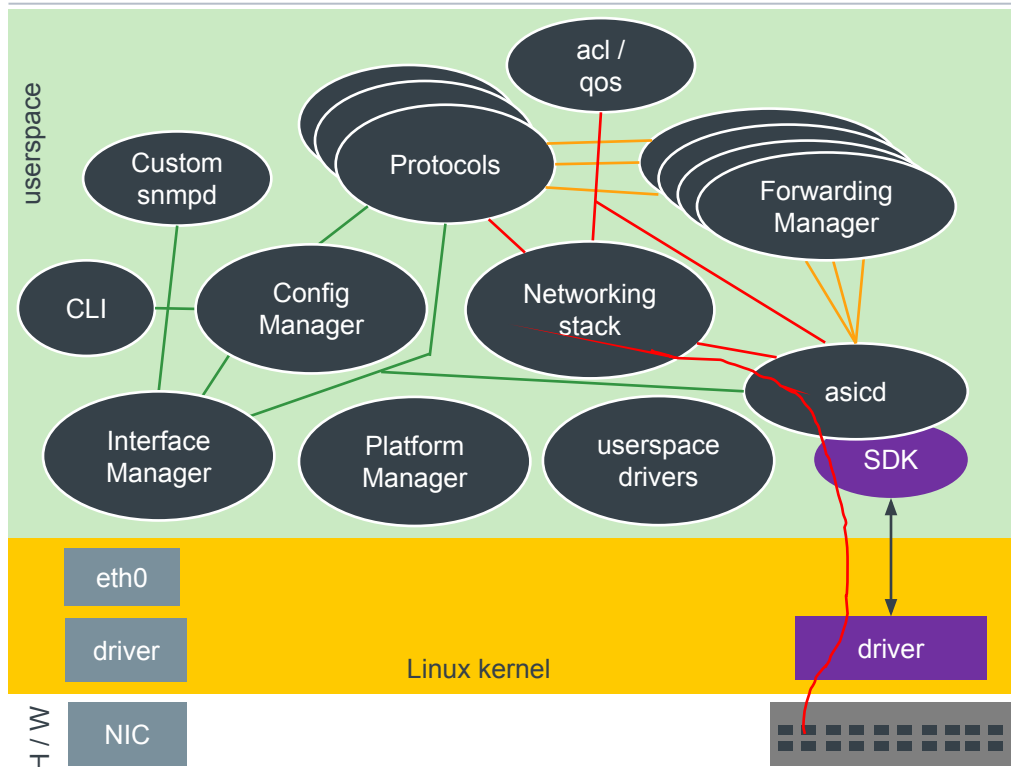
# Switch ASICs and SDKs



- Commodity ASICs controlled via SDKs
  - All commands and queries to ASIC must be go through the SDK
- NOS has userspace ASIC driver

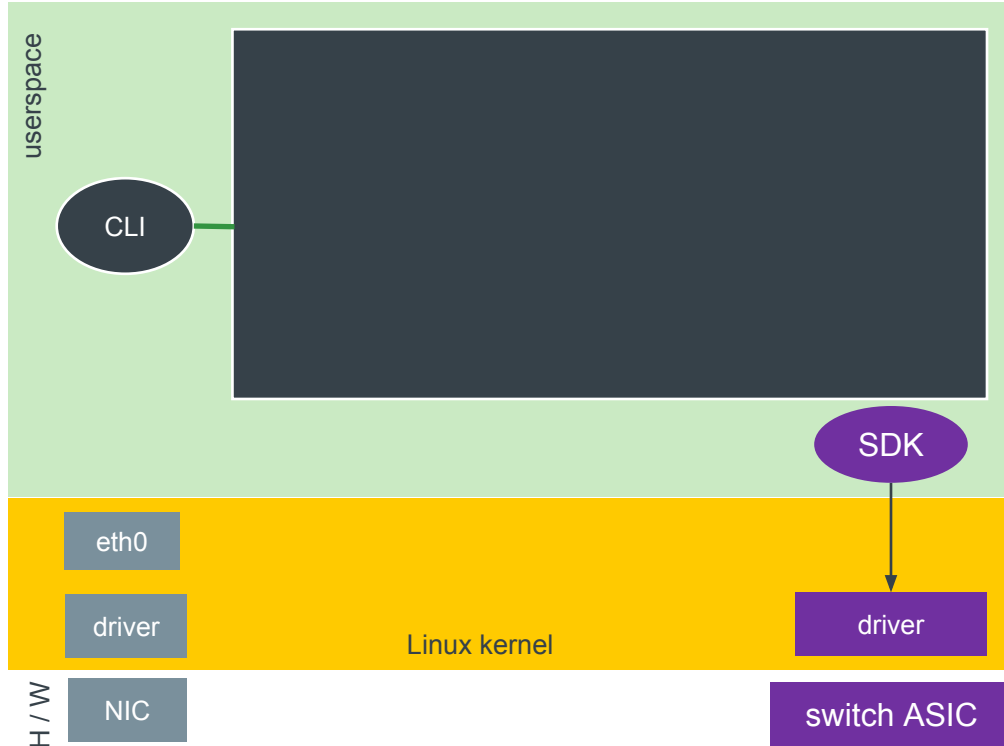


# Early NOS Architectures



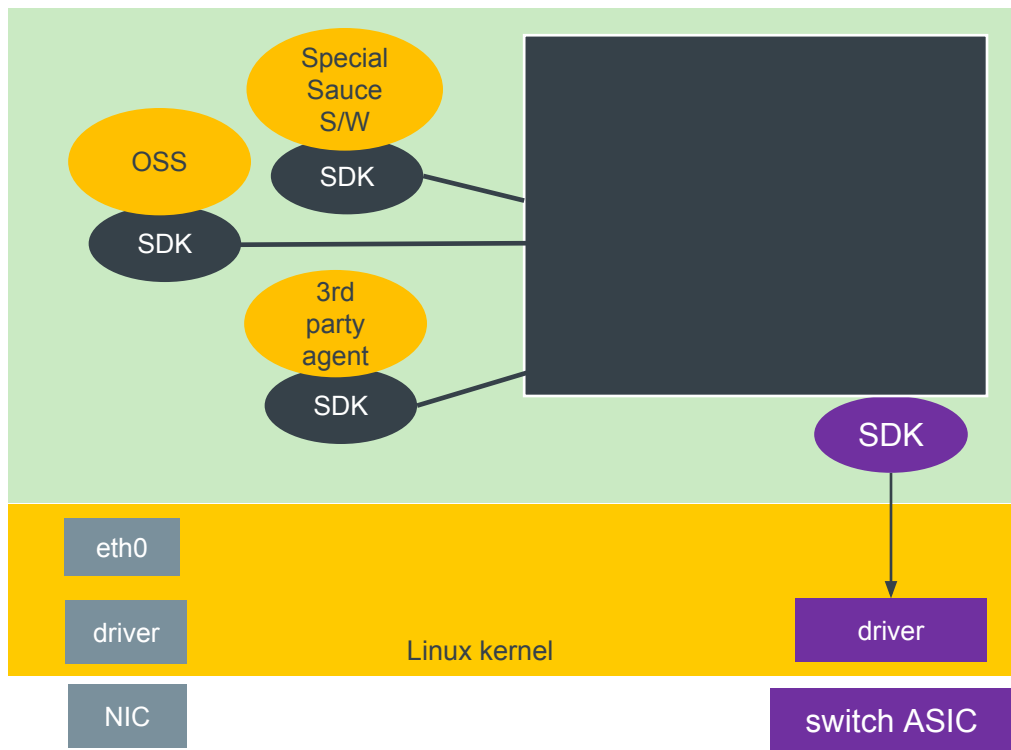
- All userspace, all custom software
  - Custom CLI to configure
  - Custom monitoring APIs
  - Custom diagnostic tools
- No data in the kernel
- Linux has a non-networking role
  - storage devices, process scheduling, management NIC, memory management, etc

# The NOS is a Black Box



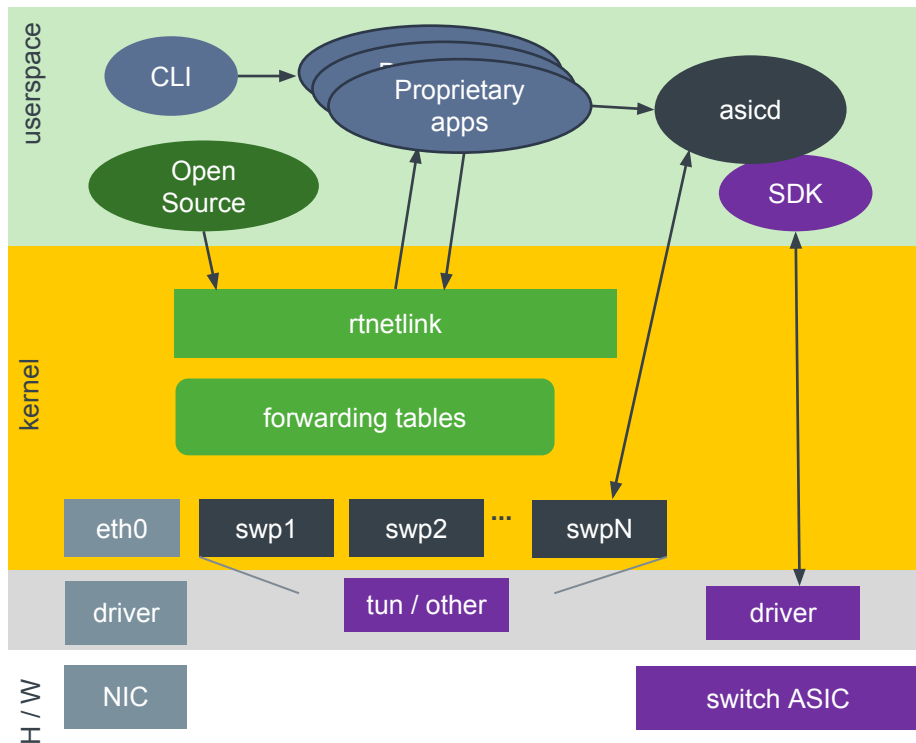
- Networking perspective it is a black box

# Legacy NOS and 3<sup>rd</sup> Party S/W



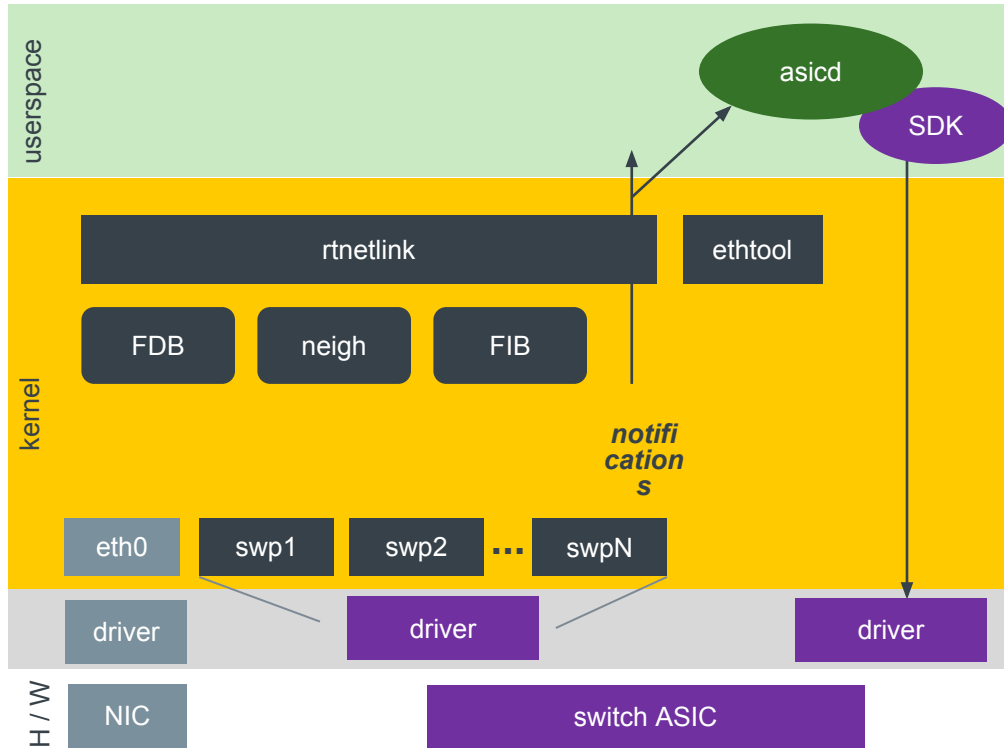
- Vendor Idea!
  - Ask customers to recompile / write their software against our SDK!
- Connect your software to the mother ship
- Seriously?

# NOS With Some Linux Networking



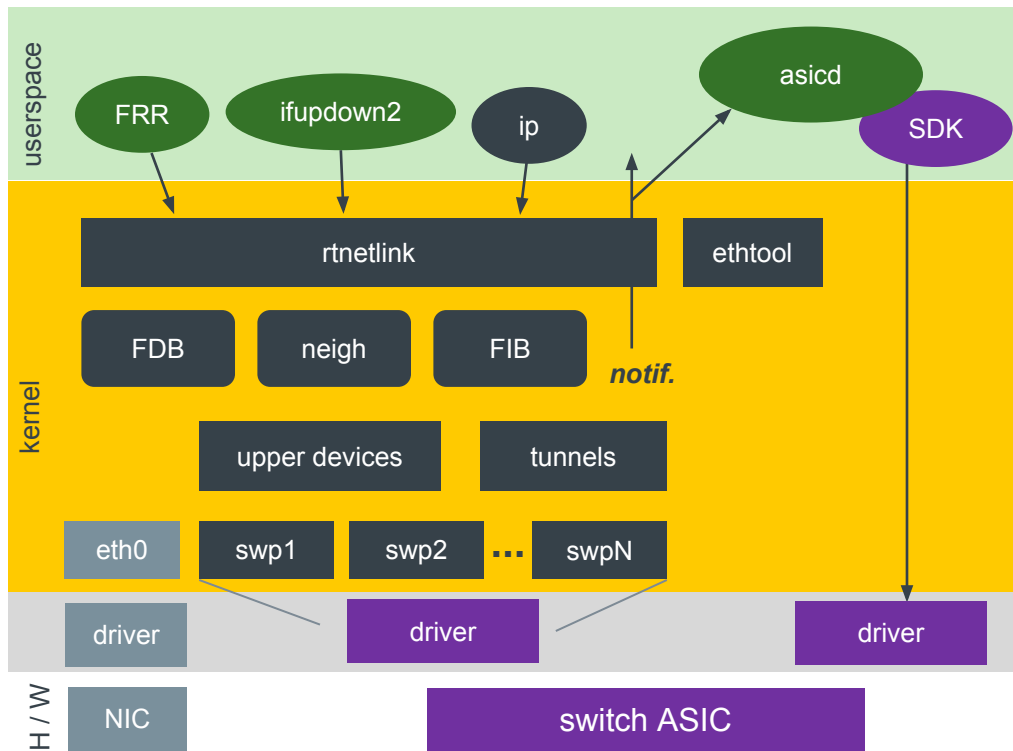
- create netdev's for ports
  - used to relay packets between SDK and Linux stack
- Pandora's box
  - How much data do you put into the kernel?
  - All of the routes? What about features - bridges, bonding, VLANs, vxlans, VRF?
  - Allow Linux APIs to configure networking, physical ports, or retrieve stats?
- Ad-hoc at best

# Adding more Linux to the equation



- Still reliant on userspace SDK
- Vendor drivers to create netdevs for front panel ports
  - in-kernel distribution of packets to port netdev
- Build from there with Linux APIs
  - rtnetlink and ethtool
  - notifications for changes to networking config and state

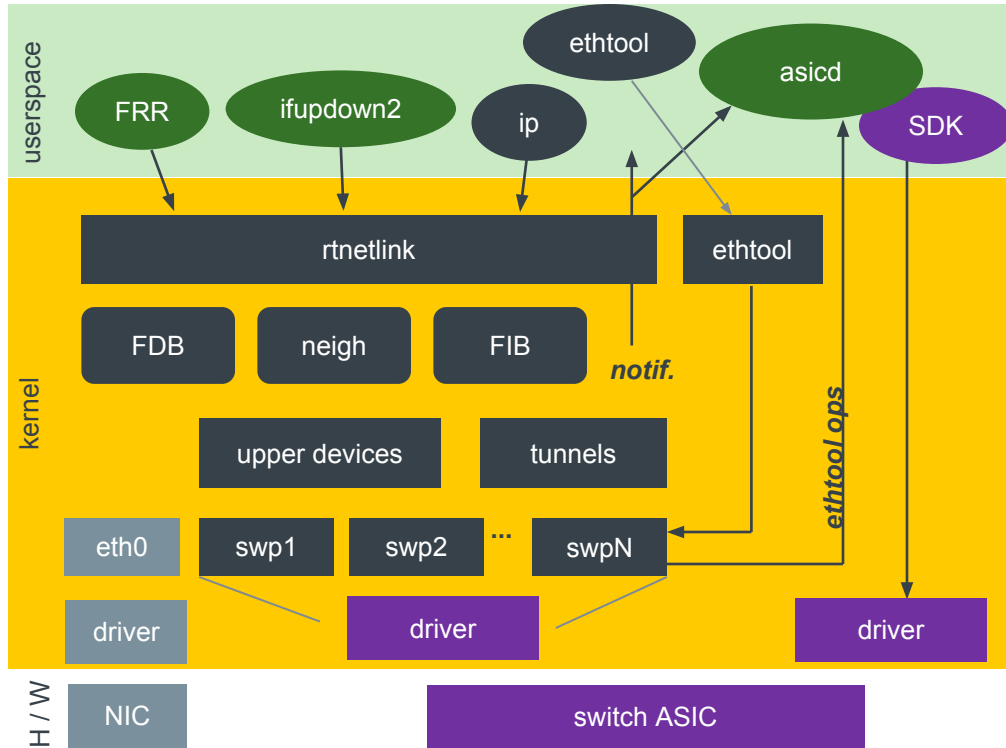
# Adding more Linux to the equation



- Now we can use the Linux ecosystem
  - Interface managers to handle complicated topologies
  - Command line tools for static / on-the-fly needs
    - e.g., `iproute2`
  - Routing suites such as FRR that speak Linux
  - Monitoring and configuration agents of your choice
    - `ansible`, `puppet`, `chef`, `collectd`, `snmpd`, ...



# Adding more Linux to the equation

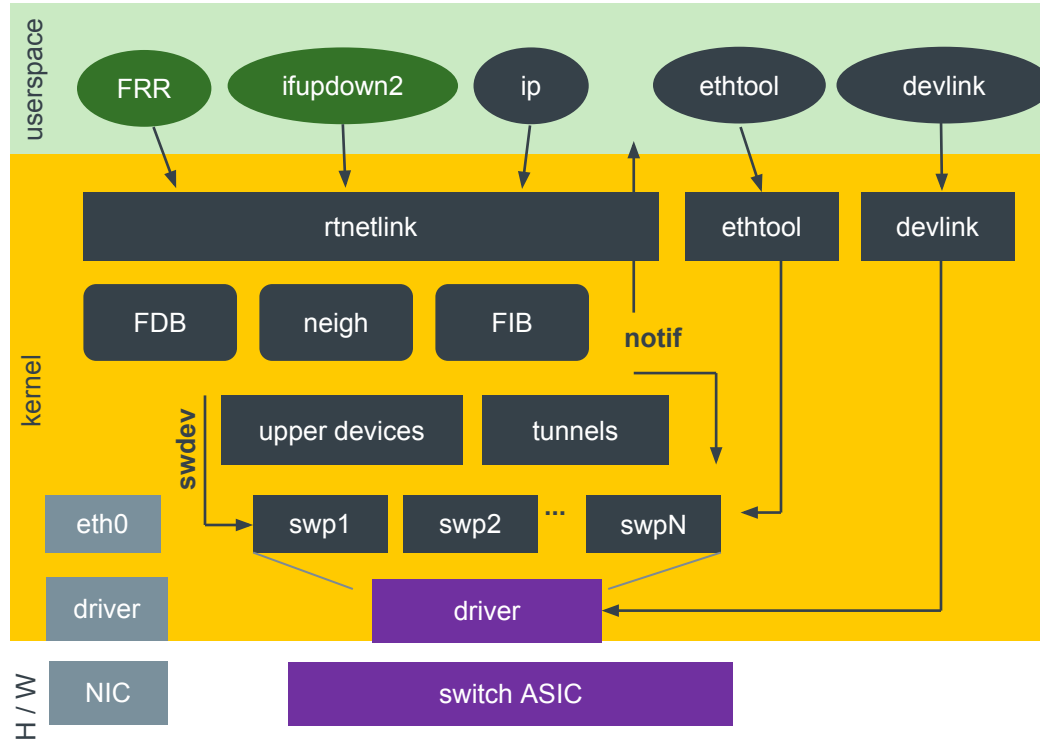


- Still a few SDK hassles
  - Configuring settings on front panel ports requires the SDK
  - Need to provide a few custom hooks
  - Error handling
  - S/W feature vs H/W feature
- Overall, much better architecture for Open Networking





# Networking offload with switchdev




- Kernel driver programs ASIC as userspace programs the kernel
- New API for users
  - devlink API for device specific data / control
- Kernel APIs for the driver to learn of changes
  - switchdev operations
  - in-kernel notifiers



# How do we get there?

---

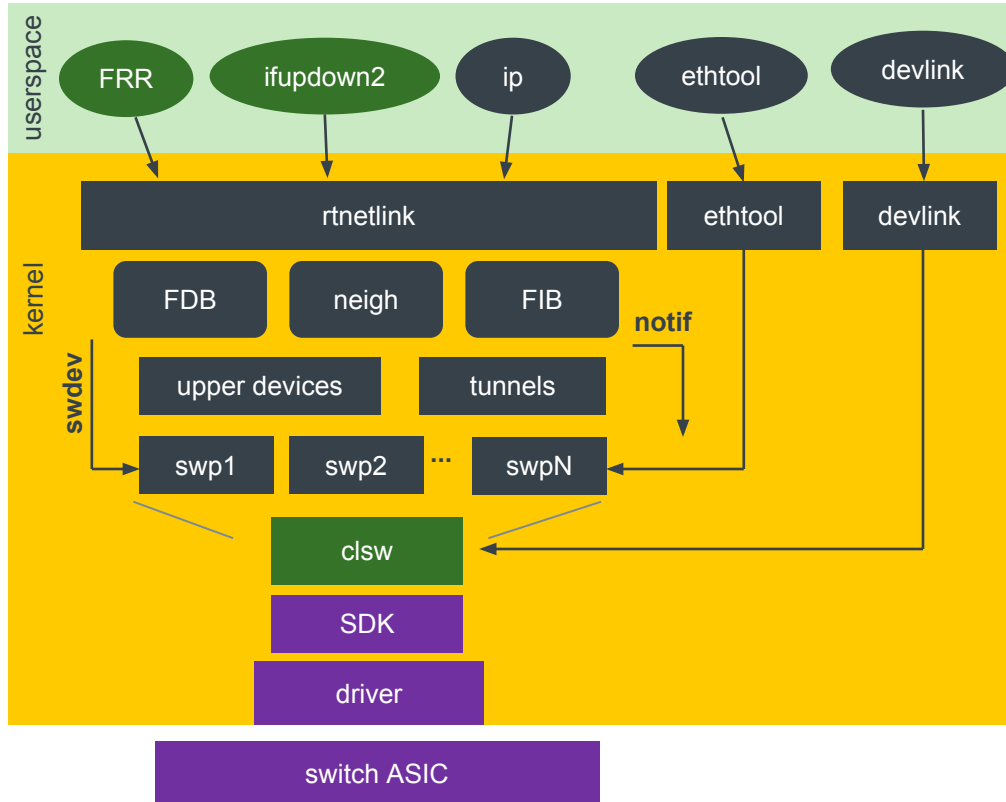
- Need ASIC vendors to support switchdev model
  - But that's a lot of work and little incentive to change from SDK
- Need to prove switchdev model is best







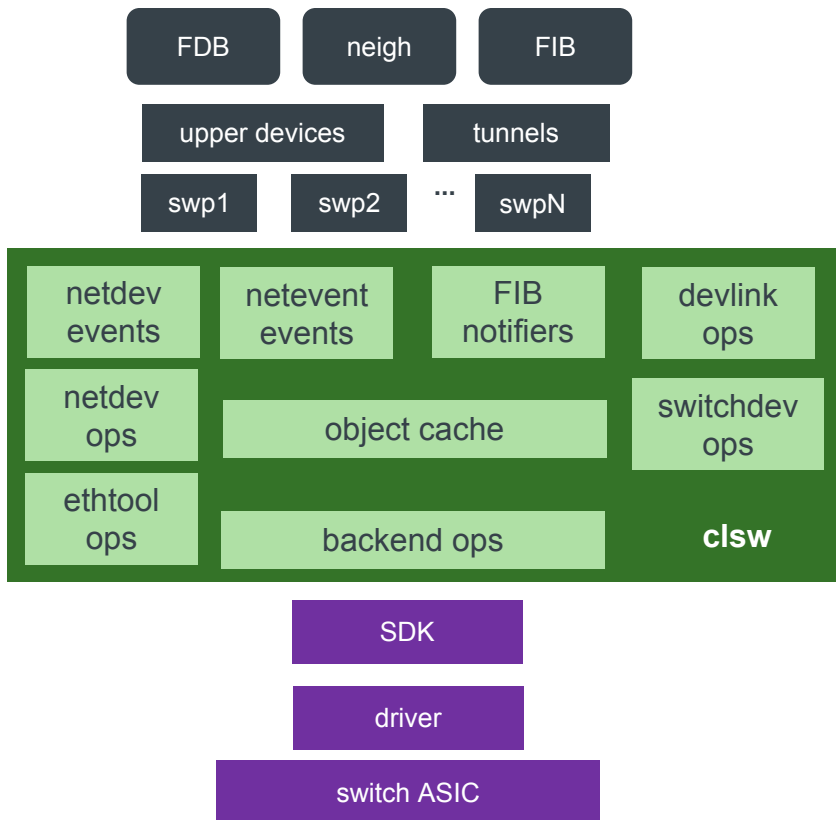
# Transitioning SDKs to the switchdev world



- Move the SDK to the kernel
- A common layer to handle kernel APIs
  - Object cache to map kernel objects to ASIC objects



# Common Layer for switchdev



- A lot of the event handling will be common across switch ASICs

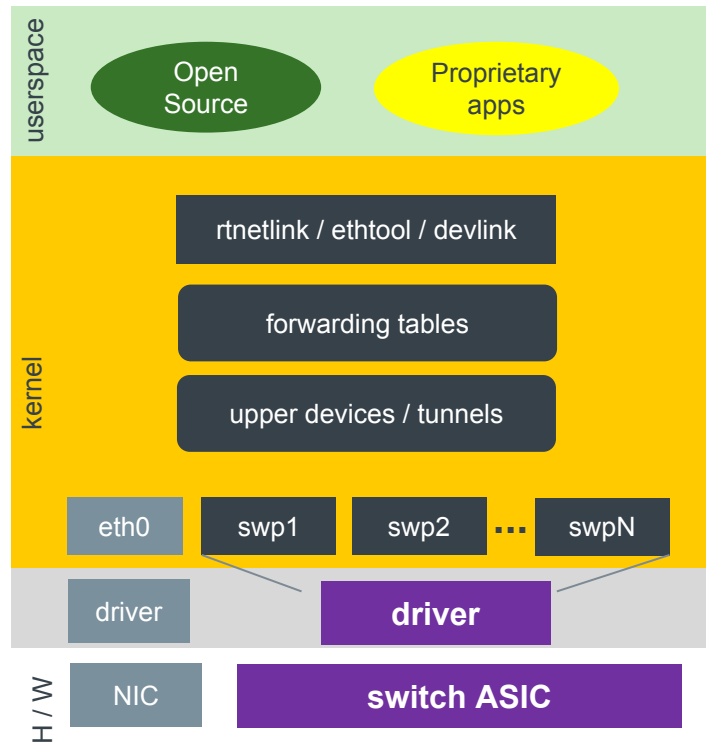
# Aren't SDKs Proprietary Blobs !?!?!





# Linux Kernel as the Center of the Open NOS


- Kernel holds networking configuration and state
  - No special process required
  - No custom API required to extract that data
- Less custom software for infrastructure pieces
  - No need for a custom IPC and the performance overhead of ping-ponging between processes
- Simpler design
  - Consistent tools / methodologies for any Linux OS
- Enables true openness and freedom
  - Use any software that speaks Linux APIs
  - Separate proprietary / unique business logic from infrastructure





# Demonstration / Proof of Concept

- Connectivity between hosts h1 and h2
  - BGP on s1 and s2, peering to exchange connected routes – establishes connectivity
- s1 is one of 3 cases:
  - sdk – userspace SDK driver
  - switchdev – 4.18-rc3 in-kernel driver
  - clsw-sdk – 4.18-rc3 in-kernel SDK based driver with clsw








## Reboot via kexec

- Number of missing ping responses between h1 and h2
  - $\text{sdk} = 72, \text{swdev} = 28, \text{clsw-sdk} = 27$
- Why?
  - Simpler initialization order
- Simpler architecture – no need for complicated “features” (ISSU) that rarely work





# Thank you!

---

Visit us at [cumulusnetworks.com](https://cumulusnetworks.com) or follow us [@cumulusnetworks](https://twitter.com/cumulusnetworks)

© 2018 Cumulus Networks. Cumulus Networks, the Cumulus Networks Logo, and Cumulus Linux are trademarks or registered trademarks of Cumulus Networks, Inc. or its affiliates in the U.S. and other countries. Other names may be trademarks of their respective owners. The registered trademark Linux® is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a world-wide basis.